

# Functions in MIPS


# Today's lecture: Implementing Functions!

- The program's flow of control must be changed.
  - The Jump and Link (jal) instruction (NEW!)
  - Using Jump Register (jr)
- Arguments and return values are passed back & forth.
  - Register Conventions
- Allocating (and deallocating) space for local variables
  - The stack
  - The stack pointer (\$sp)


# Invoking a function changes control flow by **calling** and **returning** from the function

- In this example the **main** function calls **fact** twice, and **fact** returns twice—but to *different* locations in **main**.
- Each time **fact** is called, the CPU has to remember the appropriate **return address**.
- Notice that **main** itself is also a function! It is, in effect, called by the operating system when you run the program.

```
→ int main()  
{  
    ...  
    t1 = fact(8);  
    t2 = fact(3);  
    t3 = t1 + t2;  
    ...  
}  
  
int fact(int n)  
{  
    int i, f = 1;  
    for (i = n; i > 1; i --)  
        f = f * i;  
    return f;  
}
```

A blue arrow points from the `fact(8)` call in `main` to the `fact` function definition. Another blue arrow points from the `fact(3)` call in `main` to the `fact` function definition, illustrating the return flow from the function back to the caller.

# Calling a function allocates stack frame, returning deallocates stack frame



```
int main()  
{
```

```
    ...  
    t1 = fact(8);  
    t2 = fact(3);  
    t3 = t1 + t2;
```

```
    ...  
}
```

```
int fact(int n)  
{
```

```
    int i, f = 1;  
    for (i = n; i > 1; i--)  
        f = f * i;
```

```
    return f;  
}
```

# Use `jal` to call functions, `jr` to return

- `jal` saves the return address (the address of the *next* instruction) in the dedicated register `$ra`, before jumping to the function.

3)                      PC ← Fact  
         jal Fact                      → R[31] ← PC + 4

- To transfer control back to the caller, the function just has to jump to the address that was stored in **\$ra**.

jr \$ra

# Go to Handout!

Let's call our functions and return from them

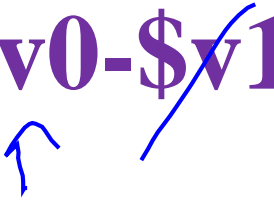
# Functions accept **arguments** and produce **return** values.

- The **blue** parts of the program show the actual and formal arguments of the fact function.
- The **purple** parts of the code deal with returning and using a result.

```
int main()
{
    ...
    t1 = fact(8);
    t2 = fact(3);
    t3 = t1 + t2;
    ...
}

int fact(int n)
{
    int i, f = 1;
    for (i = n; i > 1; i--)
        f = f * i;
    return f;
}
```

**By convention, MIPS uses \$a0-\$a3 for arguments and \$v0-\$v1 for return values**

Two hand-drawn blue arrows originate from the text. One arrow points from the register range "\$v0-\$v1" to the register "\$v0". The other arrow points from the register range "\$a0-\$a3" to the register "\$a3".

- Conventions are not enforced by the hardware or assembler, but programmers agree to them so functions written by different people can interface with each other.
- Later we'll talk about handling additional arguments or return values.



# Assembly language is **untyped**, you need to “type check your code

Untyped —there is no distinction between integers, characters, pointers or other kinds of values.

```
int x = 42  
... = fact(&x);
```

# There is a big problem here!

- The main code stores the result of fact (8) in `$t1`, but `$t1` is also used within the fact function!
- The subsequent call to fact(3) will overwrite the value of fact(8) that was stored in `$t1`.

# Calling a function within another function (**Nested functions**) can overwrite values we need

Let's say A calls B, which calls C.

- The arguments for the call to C would be placed in \$a0-\$a3, thus *overwriting* the original arguments for B.
- Similarly, **jal C** overwrites the return address that was saved in \$ra by the earlier **jal B**.

A: ...  
# Put B's args in \$a0-\$a3  
**jal B** # \$ra = A2  
A2: ...

B: ...  
# Put C's args in \$a0-\$a3,  
# erasing B's args!  
**jal C** # \$ra = **B2**  
B2: ...  
~~j r \$ra~~ # Where does  
# this go???

C: ...  
j r \$ra

- a) A
- b) A2
- c) B
- d) B2
- e) C

# Spilling registers



- The CPU has a limited number of registers for use by all functions, and it's possible that several functions will need the same registers.
- We can keep important registers from being overwritten by a function call, by saving them before the function executes, and restoring them after the function completes.
- But there are two important questions.
  - Who is responsible for saving registers—the caller or the callee?
  - Where exactly are the register contents saved?

# Who saves the registers?

- Who is responsible for saving important registers across function calls?
  - The caller knows which registers are important to it and should be saved.
  - The callee knows exactly which registers it will use and potentially overwrite.
- However, in the typical “black box” programming approach, the caller and callee do not know anything about each other’s implementation.
  - Different functions may be written by different people or companies.
  - A function should be able to interface with any client, and different implementations of the same function should be substitutable.
- So how can two functions cooperate and share registers when they don’t know anything about each other?

# The caller could save the registers...

- One possibility is for the *caller* to save any important registers that it needs before making a function call, and to restore them after.
- But the caller does not know what registers are actually written by the function, so it may save more registers than necessary.
- In the example on the right, **frodo** wants to preserve **\$a0**, **\$a1**, **\$s0** and **\$s1** from **gollum**, but gollum may not even use those registers.

```
frodo: li    $a0, 3  
       li    $a1, 1  
       li    $s0, 4  
       li    $s1, 1
```

```
# Save registers  
# $a0, $a1, $s0, $s1
```

```
j al gollum
```



```
# Restore registers  
# $a0, $a1, $s0, $s1
```

```
add $v0, $a0, $a1  
add $v1, $s0, $s1  
jr   $ra
```


# ...or the callee could save the registers...

- Another possibility is if the *callee* saves and restores any registers it might overwrite.
- For instance, a `gollum` function that uses registers `$a0`, `$a2`, `$s0` and `$s2` could save the original values first, and restore them before returning.
- But the callee does not know what registers are important to the caller, so again it may save more registers than necessary.

`gollum:`

```
# Save registers  
# $a0 $a2 $s0 $s2
```

```
li    $a0, 2  
li    $a2, 7  
li    $s0, 1  
li    $s2, 8  
...
```



```
# Restore registers  
# $a0 $a2 $s0 $s2
```

```
jr    $ra
```

# ...or they could work together

- MIPS uses conventions again to split the register spilling chores.
- The *caller* is responsible for saving and restoring any of the following **caller-saved registers** that it cares about.

\$t0-\$t9

\$a0-\$a3

\$v0-\$v1

In other words, the callee may freely modify these registers, under the assumption that the caller already saved them if necessary.

- The *callee* is responsible for saving and restoring any of the following **callee-saved registers** that it uses.

\$s0-\$s7

Thus the caller may assume these registers are not changed by the callee.

- ~~\$ra~~ is special; it is “used” by jal. It is saved by a callee who is also a caller.

\$ra



# Register spilling example

This convention ensures that the caller and callee together save all of the important registers—frodo only needs to save registers `$a0` and `$a1`, while gollum only has to save registers `$s0` and `$s2`.

```
frodo: li    $a0, 3
       li    $a1, 1
       li    $s0, 4
       li    $s1, 1

       # Save registers
       # $a0, $a1, $ra

       jal   gollum

       # Restore registers
       # $a0, $a1, $ra

       add   $v0, $a0, $a1
       add   $v1, $s0, $s1
       jr    $ra
```

```
gollum:
       # Save registers
       # $s0 and $s2

       li    $a0, 2
       li    $a2, 7
       li    $s0, 1
       li    $s2, 8
       ...

       # Restore registers
       # $s0 and $s2


       jr    $ra
```

# In the factorial example, main (the caller) should save two registers

- `$t1` must be saved before the second call to fact.
- `$ra` will be implicitly overwritten by the jal instructions.
- But fact (the callee) does not need to save anything. It only writes to registers `$t0`, `$t1` and `$v0`, which should have been saved by the caller.

# Where are the registers saved?

- Now we know who is responsible for saving which registers, but we still need to discuss where those registers are saved.
- It would be nice if each function call had its own private memory area.
  - This would prevent other function calls from overwriting our saved registers—otherwise using memory is no better than using registers.
  - We could use this private memory for other purposes too, like storing local variables.




# Use the stack for caller and callee saves

- Notice function calls and returns occur in a stack-like order: the most recently called function is the first one to return.

1. Someone calls A
2. A calls B
3. B calls C
4. C returns to B
5. B returns to A
6. A returns

- Here, for example, C must return to B *before* B can return to A.




# Stacks and function calls

- It's natural to use a **stack** for function call storage. A block of stack space, called a **stack frame**, can be allocated for each function call.
  - When a function is called, it creates a new frame onto the stack, which will be used for local storage.
  - Before the function returns, it must pop its stack frame, to restore the stack to its original state.
- The stack frame can be used for several purposes.
  - Caller- and callee-save registers can be put in the stack.
  - The stack frame can also hold local variables, or extra arguments and return values.



# The MIPS stack

- In MIPS machines, part of main memory is reserved for a stack.
  - The stack grows downward in terms of memory addresses.
  - The address of the top element of the stack is stored (by convention) in the “stack pointer” register, **\$sp**.
- MIPS does not provide “push” and “pop” instructions. Instead, they must be done explicitly by the programmer.



# Pushing elements

- To **push** elements onto the stack:
  - Move the stack pointer **\$sp** down to make room for the new data.
  - Store the elements into the stack.
- For example, to push registers **\$t1** and **\$t2** onto the stack:

```
sub $sp, $sp, 8  
sw  $t1, 4($sp)  
sw  $t2, 0($sp)
```



Handwritten annotations: a blue bracket on the right side of the code block, and a blue arrow pointing to the `0($sp)` offset.

- An equivalent sequence is:

```
sw  $t1, -4($sp)  
sw  $t2, -8($sp)  
sub $sp, $sp, 8
```

Handwritten annotations: a large blue 'X' to the right of the code block.

- Before and after diagrams of the stack are shown on the right.



# Accessing and popping elements

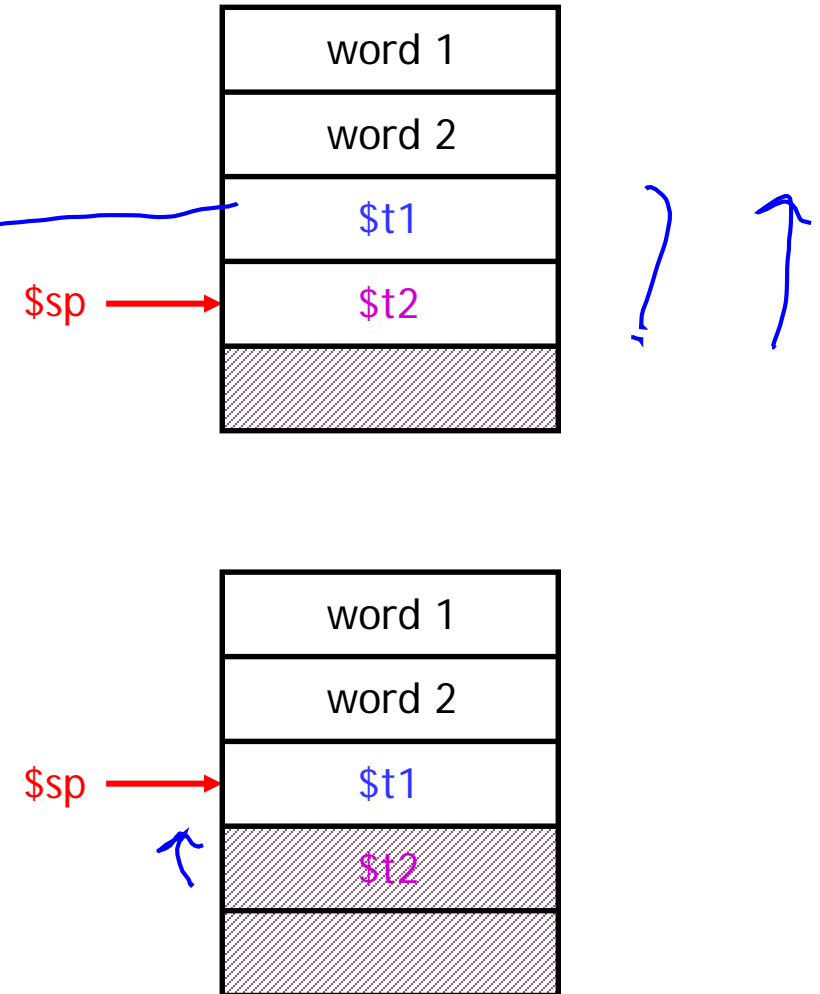
- You can access any element in the stack (not just the top one) if you know where it is relative to `$sp`.
- For example, to retrieve the value of `$t1`:

`lw $s0, 4($sp)`

- You can **pop**, or “erase,” elements simply by adjusting the stack pointer upwards.
- To pop the value of `$t2`, yielding the stack shown at the bottom:

`addi $sp, $sp, 4`

- Note that the popped data is still present in memory, but data past the stack pointer is considered invalid.





# Summary

- Today we focused on implementing function calls in MIPS.
  - We call functions using `jal`, passing arguments in registers `$a0-$a3`.
  - Functions place results in `$v0-$v1` and return using `jr $ra`.
- Managing resources is an important part of function calls.
  - To keep important data from being overwritten, registers are saved according to conventions for `caller-save` and `callee-save` registers.
  - Each function call uses stack memory for saving registers, storing local variables and passing extra arguments and return values.
- Assembly programmers must follow many conventions. Nothing prevents a rogue program from overwriting registers or stack memory used by some other function.
- On Monday, we'll look at writing recursive functions.