

Stalls and Flushes

MONDAY: Stalls and Flushes


Wednesday: Pipelining since 1985 (optional)

Friday: MIPS coding review session (optional)

Today's Lecture

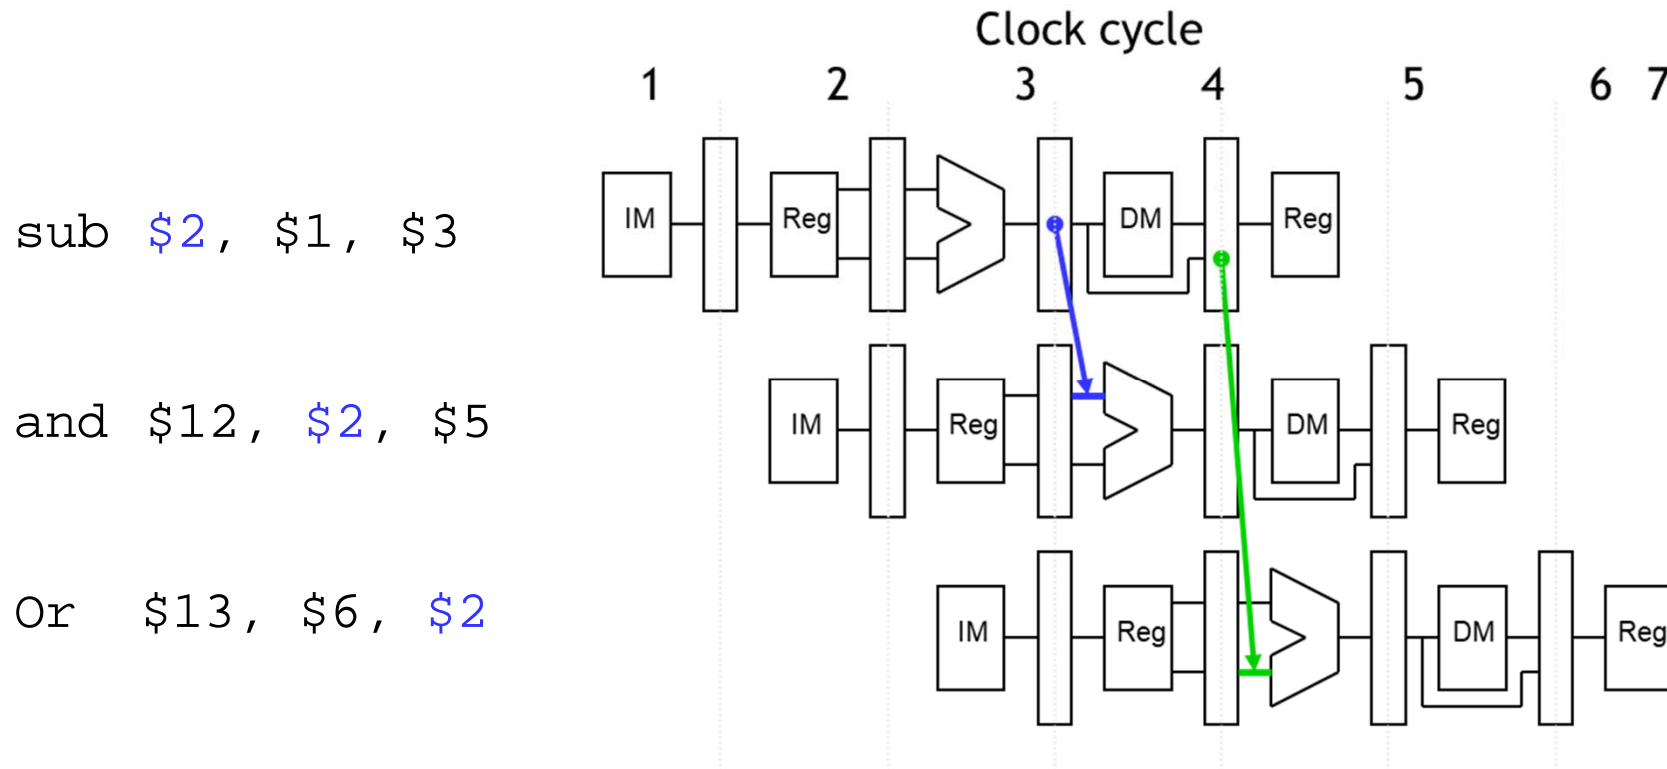
- Stalls
 - Some data dependencies that cannot be resolved with only forwarding
 - Example: Loads
 - Example: Branches
- Flushes
 - Some data dependencies cause us to fetch the wrong instructions
 - Flushes remove wrong instructions from the pipeline

Use arrows to show dependencies: Arrows that point backwards reveal **data hazards**



Forward values from pipeline registers so later instructions can use the correct value

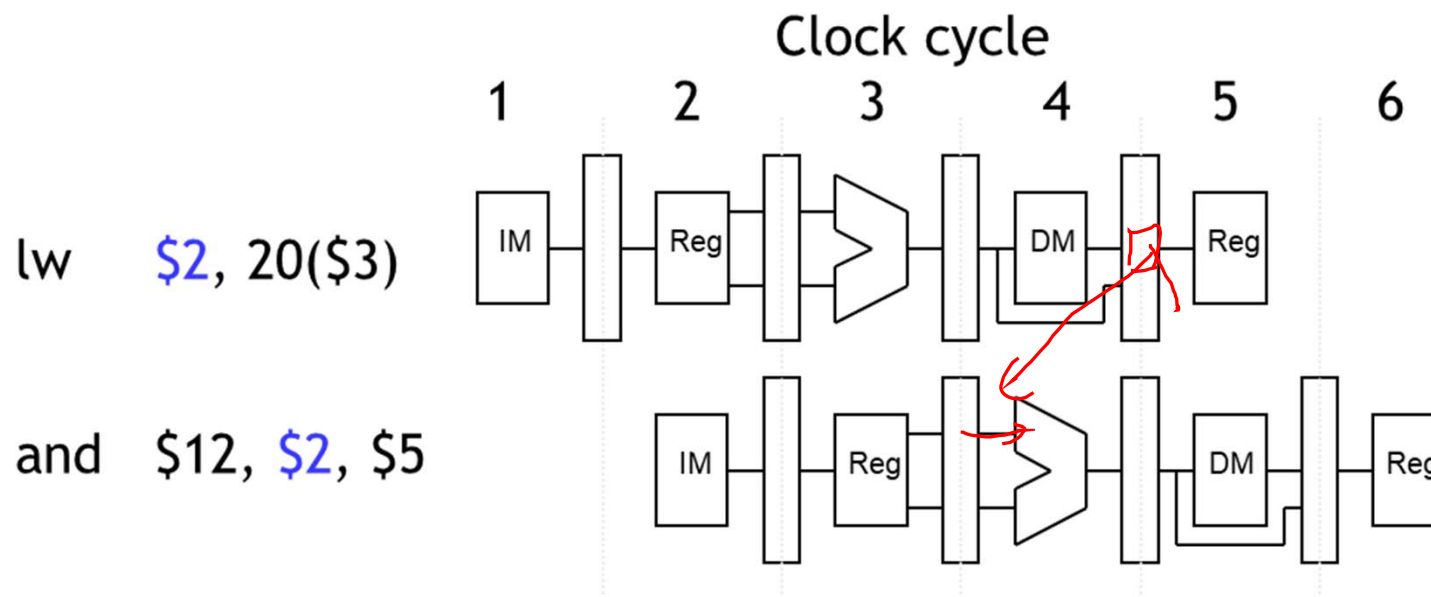
- In clock cycle 4, AND gets R[1]-R[3] from EX/MEM
- In cycle 5, OR gets R[1]-R[3] from MEM/WB



Consider the following sequence of instructions

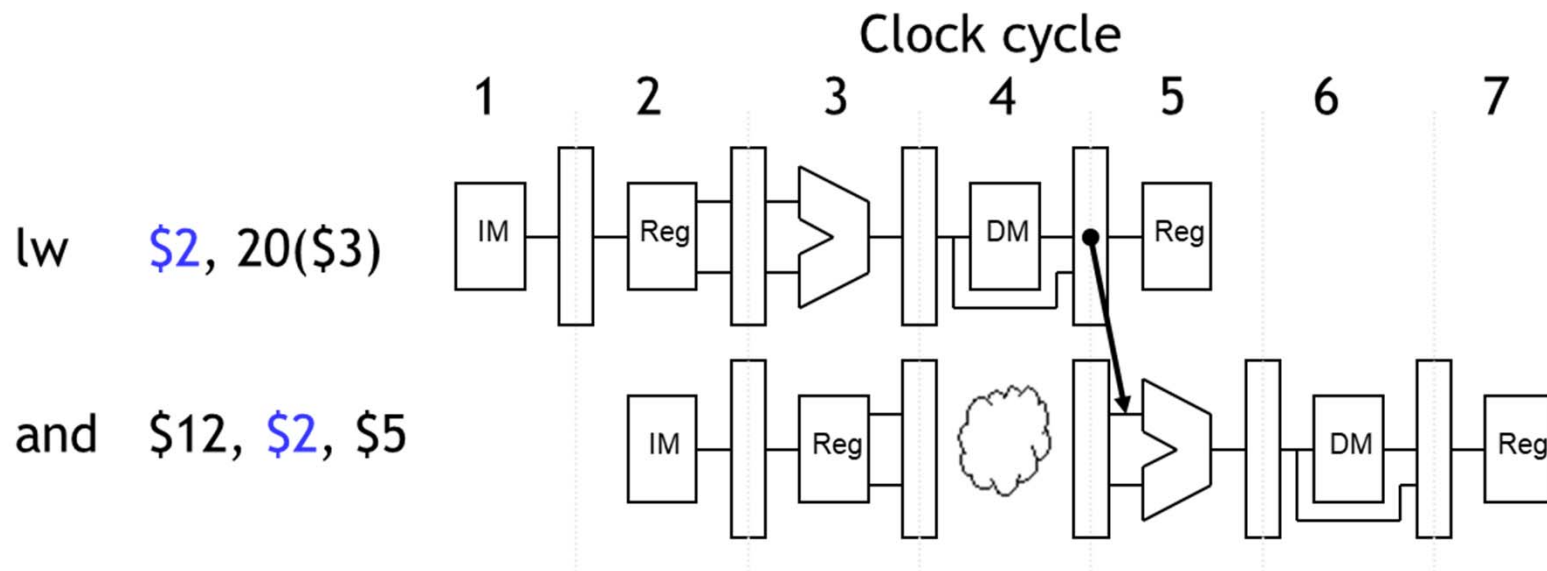
Can we resolve this data hazard by forwarding data from...

- a) LW execute (cycle 3) to AND execute (cycle 4)
- b) LW memory (cycle 4) to AND execute (cycle 4)
- c) LW memory (cycle 4) to AND memory (cycle 5)
- d) None of the above



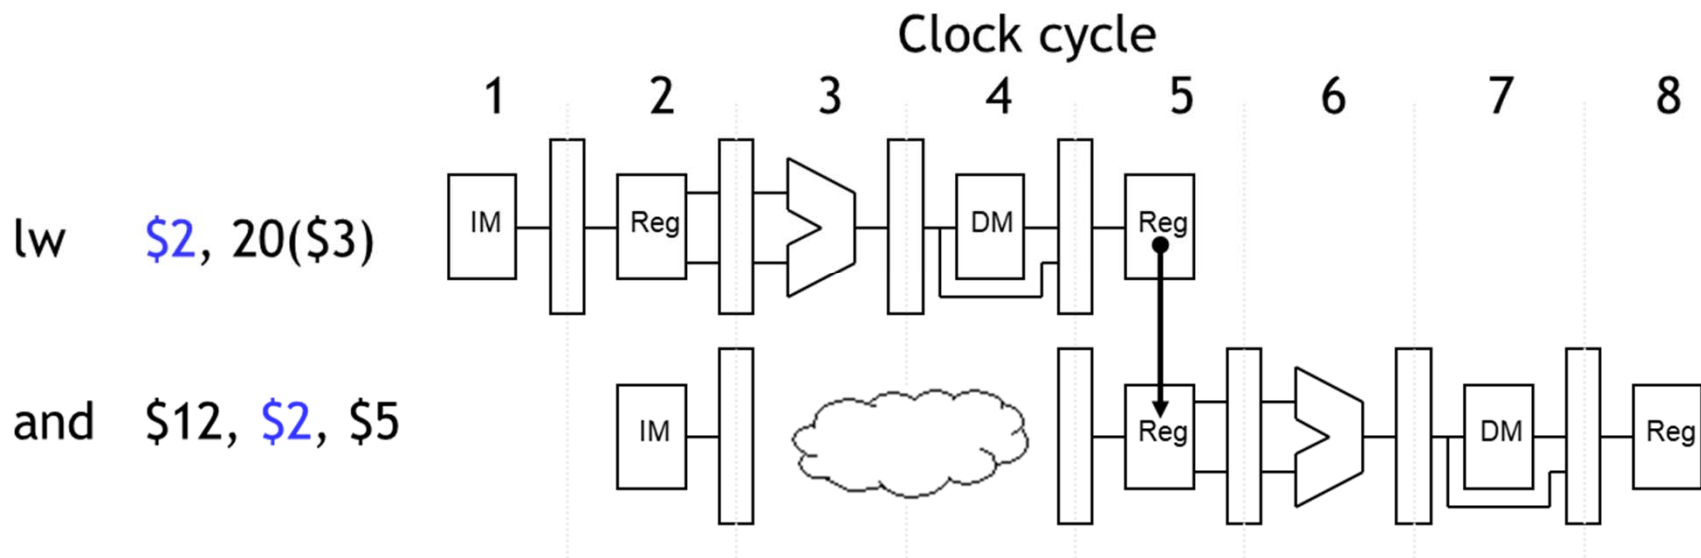
When we cannot resolve a hazard with forwarding alone, we need to **stall** the pipeline

- To stall the pipeline, introduce a one-cycle delay into the pipeline, (a **bubble**).



- Notice that we're still using forwarding in cycle 5, to get data from the MEM/WB pipeline register to the ALU.

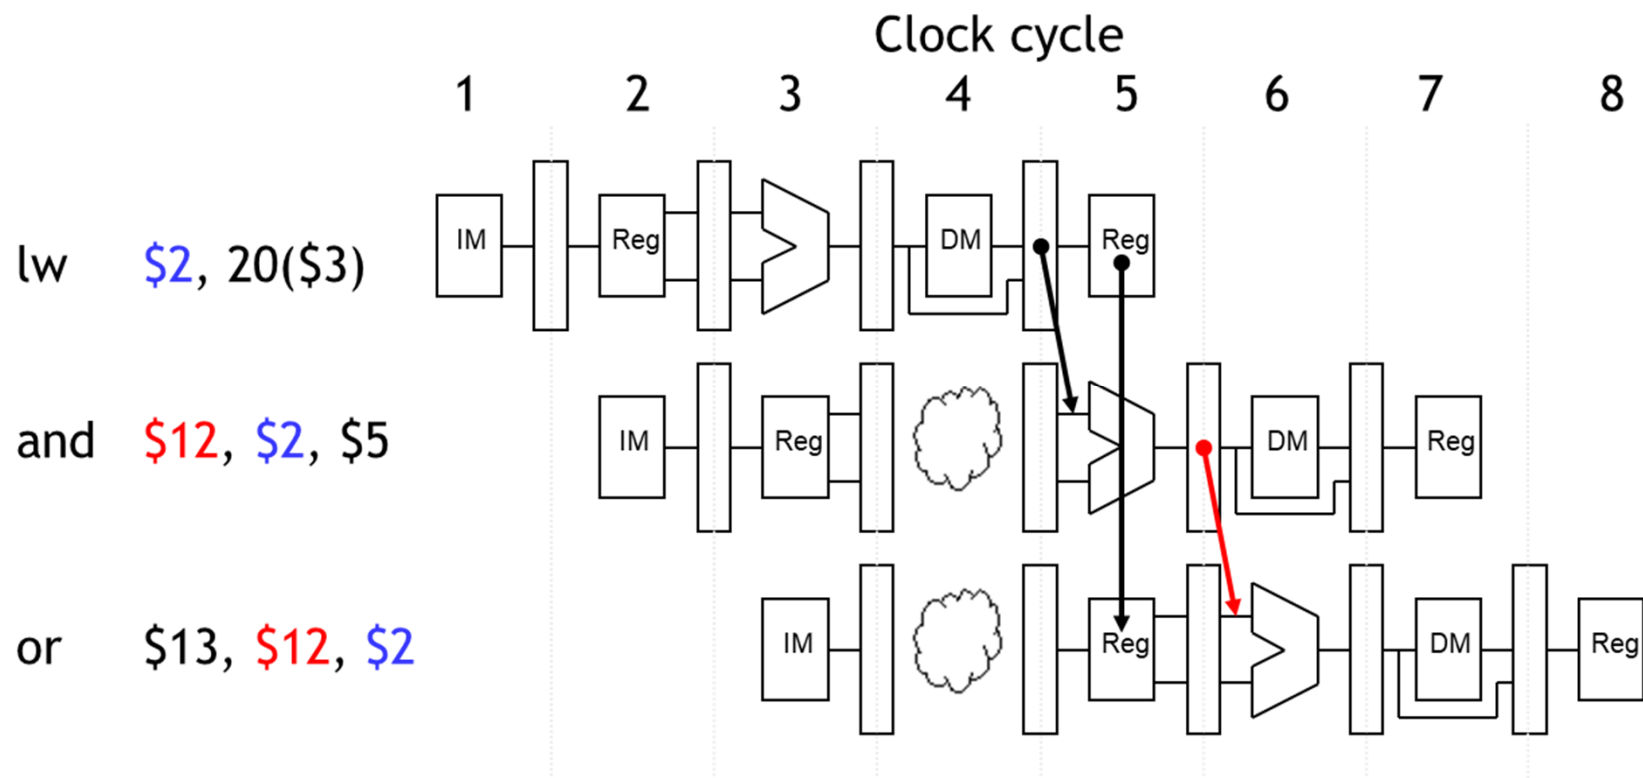
Without forwarding we would have to stall for two cycles for `lw` to reach writeback



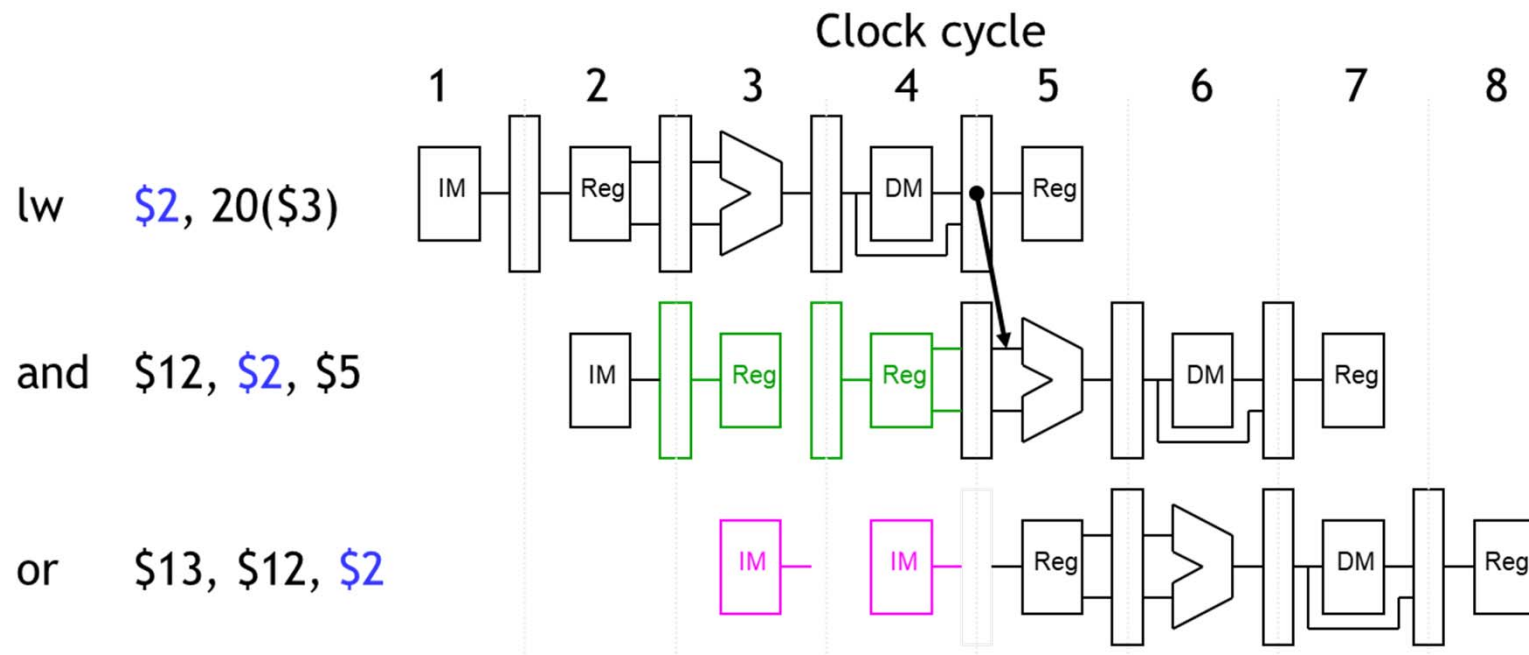
- In general, you can always stall to avoid hazards—but dependencies are very common in real code, and stalling often can reduce performance by a significant amount.

Stalling delays the entire pipeline

- Delay the second instruction -> delay the third instruction too.
 - Why? (two reasons)

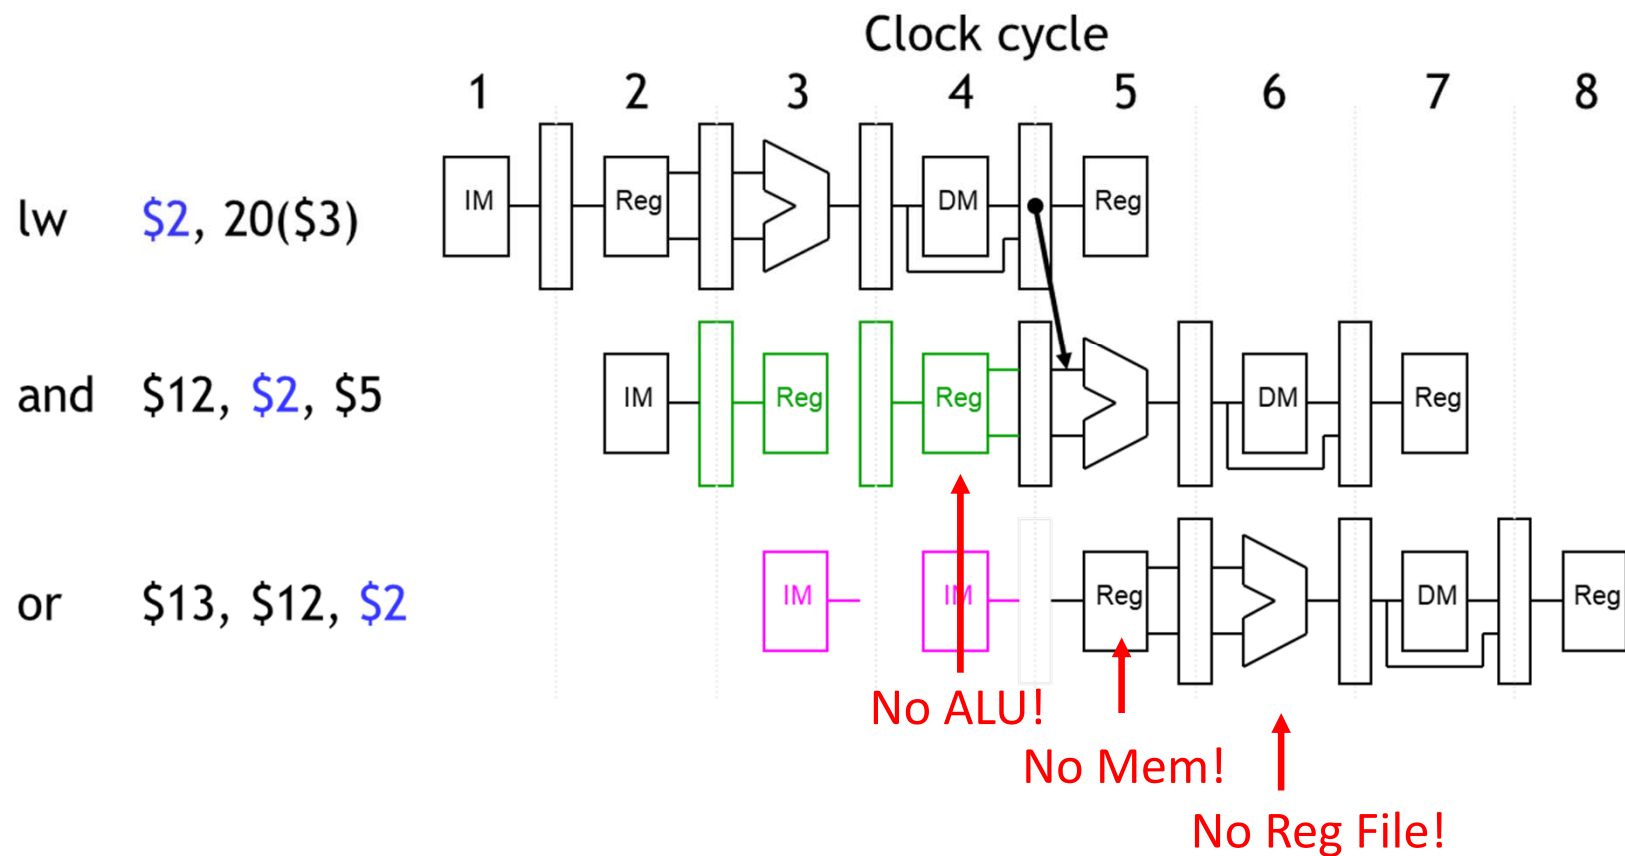


Step 1: Force the two instructions after `lw` to remain in ID and IF stages for one extra cycle




- This is easily accomplished.
 - Don't update the PC, so the current IF stage is repeated.
 - Don't update the IF/ID register, so the ID stage is also repeated.

Step 2: Set control signals for EX, MEM, and WB to zero for cycles 4, 5, and 6 respectively



Stalling converts some instructions to nop (“no operation”) instructions



Implement the `nop` in MIPS with `sll $0, $0, 0`

Name,	Mnemonic	Format	Operation	Opcode/Funct (Hex)
Shift Left Logical	<code>sll</code>	R	$R[rd] = R[rt] \ll \underline{\text{shamt}}$	<u>0/00</u>

opcode	rs	rt	rd	shamt	funct
000000	00000	00000	00000	00000	000000

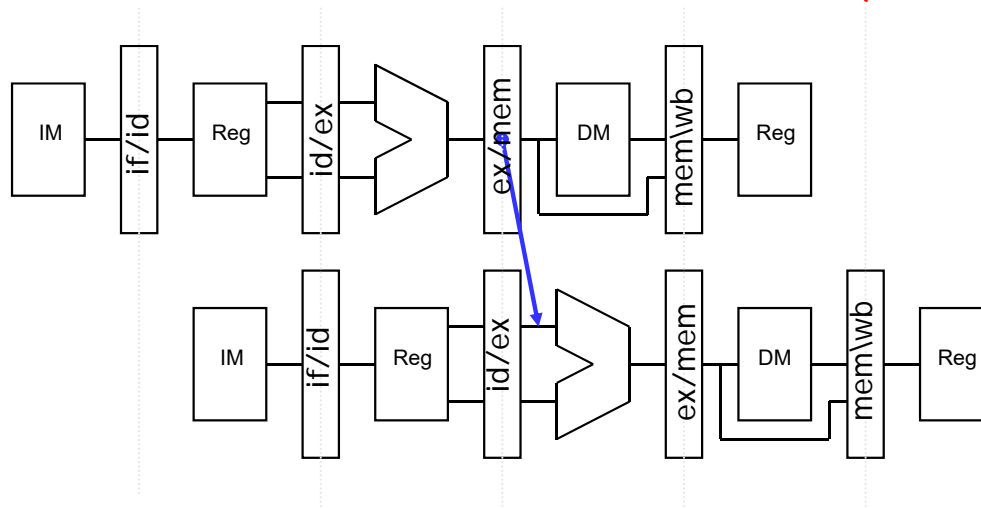
To detect stalls, we need to find dependencies

- Recall how we detected hazards that required forwarding?


if ($\text{EX/MEM.RegWrite} = 1$ *- sub match*
and $\text{EX/MEM.RegisterRd} = \text{ID/EX.RegisterRs}$ *- \$2 match*)
then Bypass Rs from EX/MEM stage latch *- forward*

sub *rd* (\$2), \$1, \$3

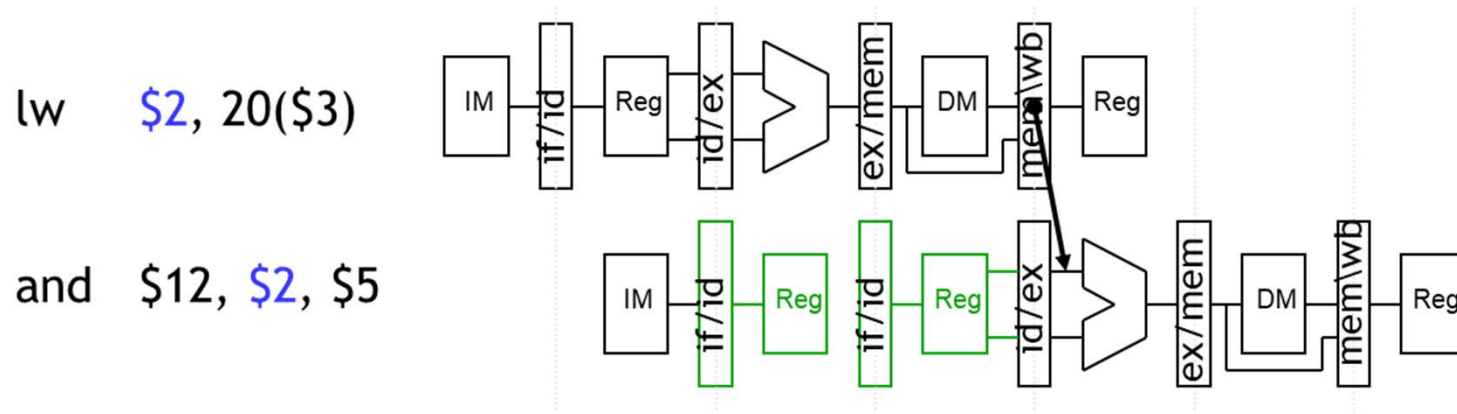
and \$12, *rs* \$2, \$5



Which set of pipeline registers should we use to detect that we need to stall?



Detect stalls when `lw` finishes the decode stage

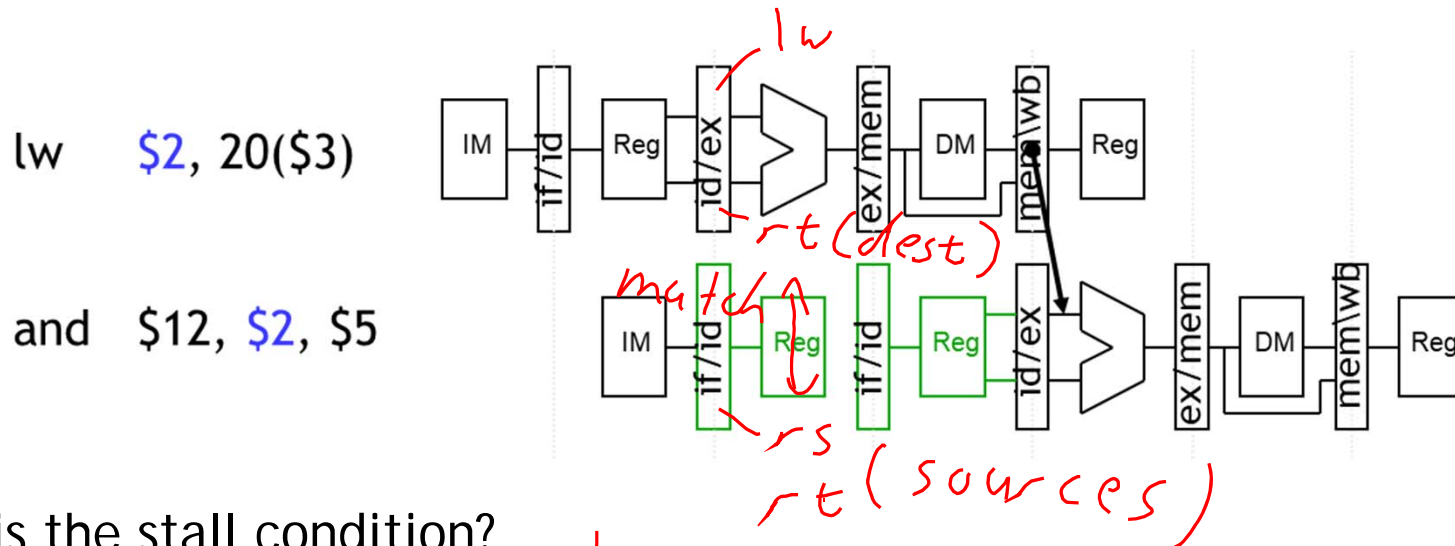


What is the stall condition?

if (

then `stall`

Detect stalls when `lw` finishes the decode stage



What is the stall condition?

if (ID/EX.MemRead = 1 and

(ID/EX.RegisterRt = IF/ID.RegisterRs or

ID/EX.RegisterRt = IF/ID.RegisterRt))

then stall

Table method for tracking stalls

clock cycles

4 filling

3 WB

1 stall

8

lw


and

or

\$2, 20(\$3)

\$12, \$2, \$5

\$13, \$12, \$2



Clock Cycle

1

2

3

4

5

1

2

3

4

5

6

7

8

9


F	D	E	M	W				
	F	D	D	E	M	W		
		F	F	D	E	M	W	

1 stall

3 WB

Table method for tracking stalls

There
is
forwarding



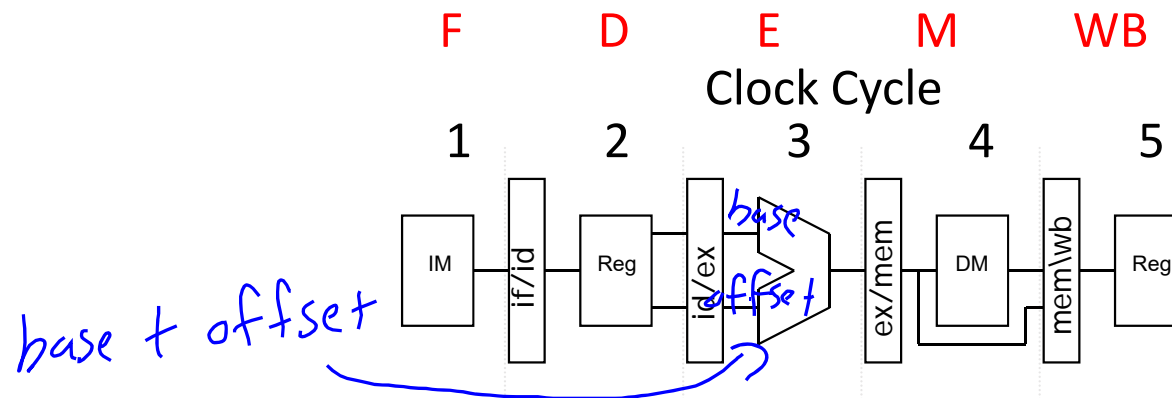
lw \$9, 4(\$7)
and \$2, \$7, \$8
add \$4, \$9, \$2

	1	2	3	4	5	6	7	8	9
lw	F	D	E	M	W				
and		F	D	E	M	W			
add			F	D	E	M	W		
				A	B	C	D	E	

In which cycle is
add \$4, \$9, \$2
in the WB stage?

+

Table method for tracking stalls

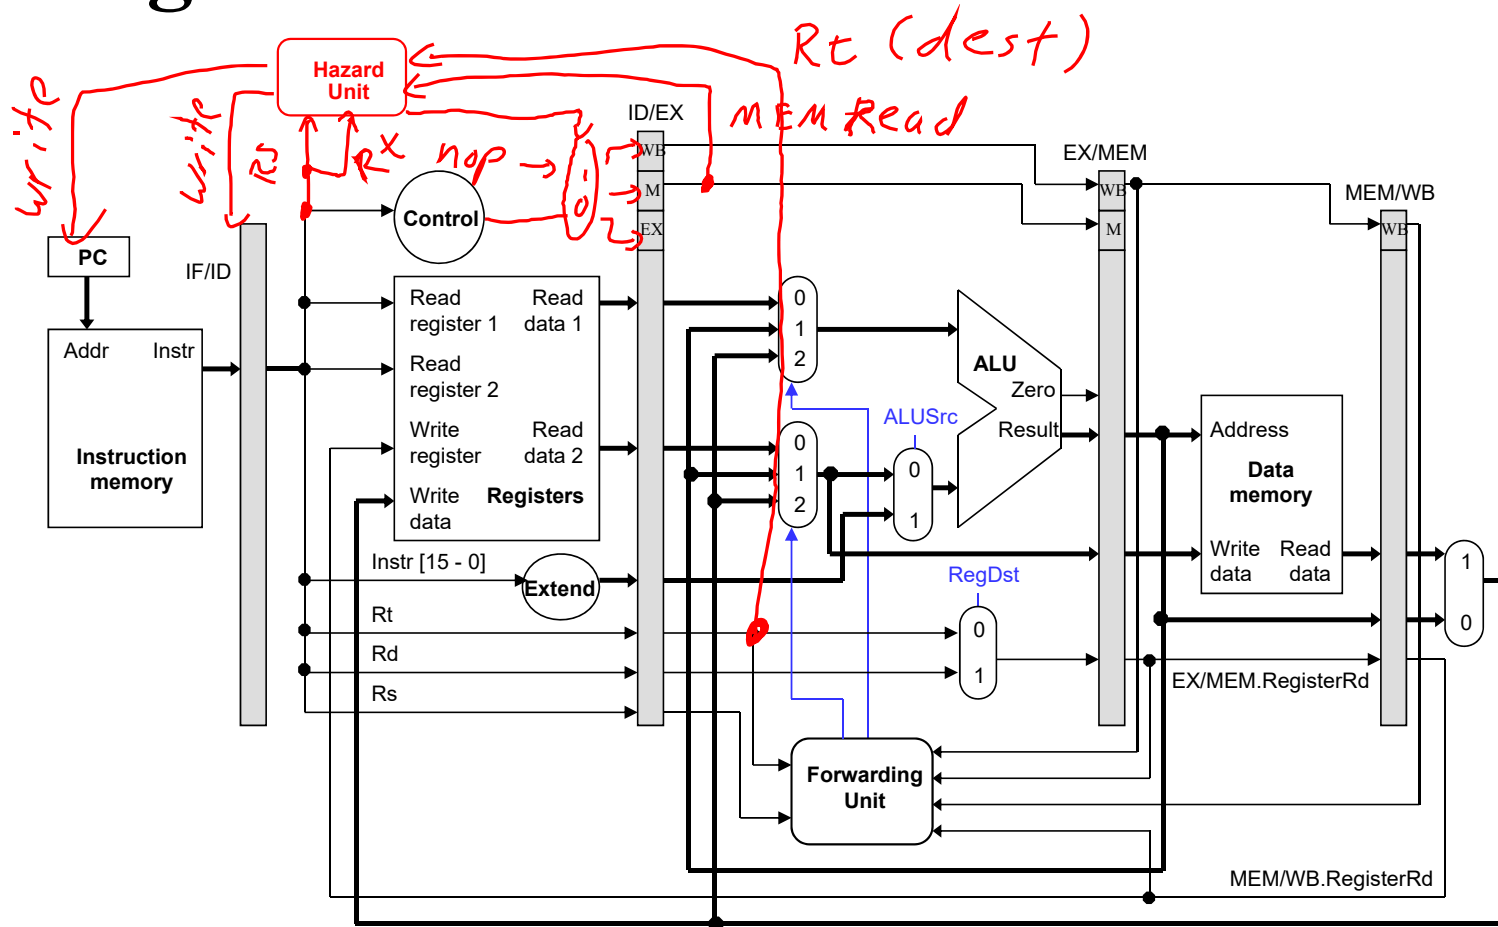


lw (\$9), 4(\$7)
 lw (\$5), 8(\$9)
 and (\$2), \$5, \$8
 add \$4, \$9, \$2

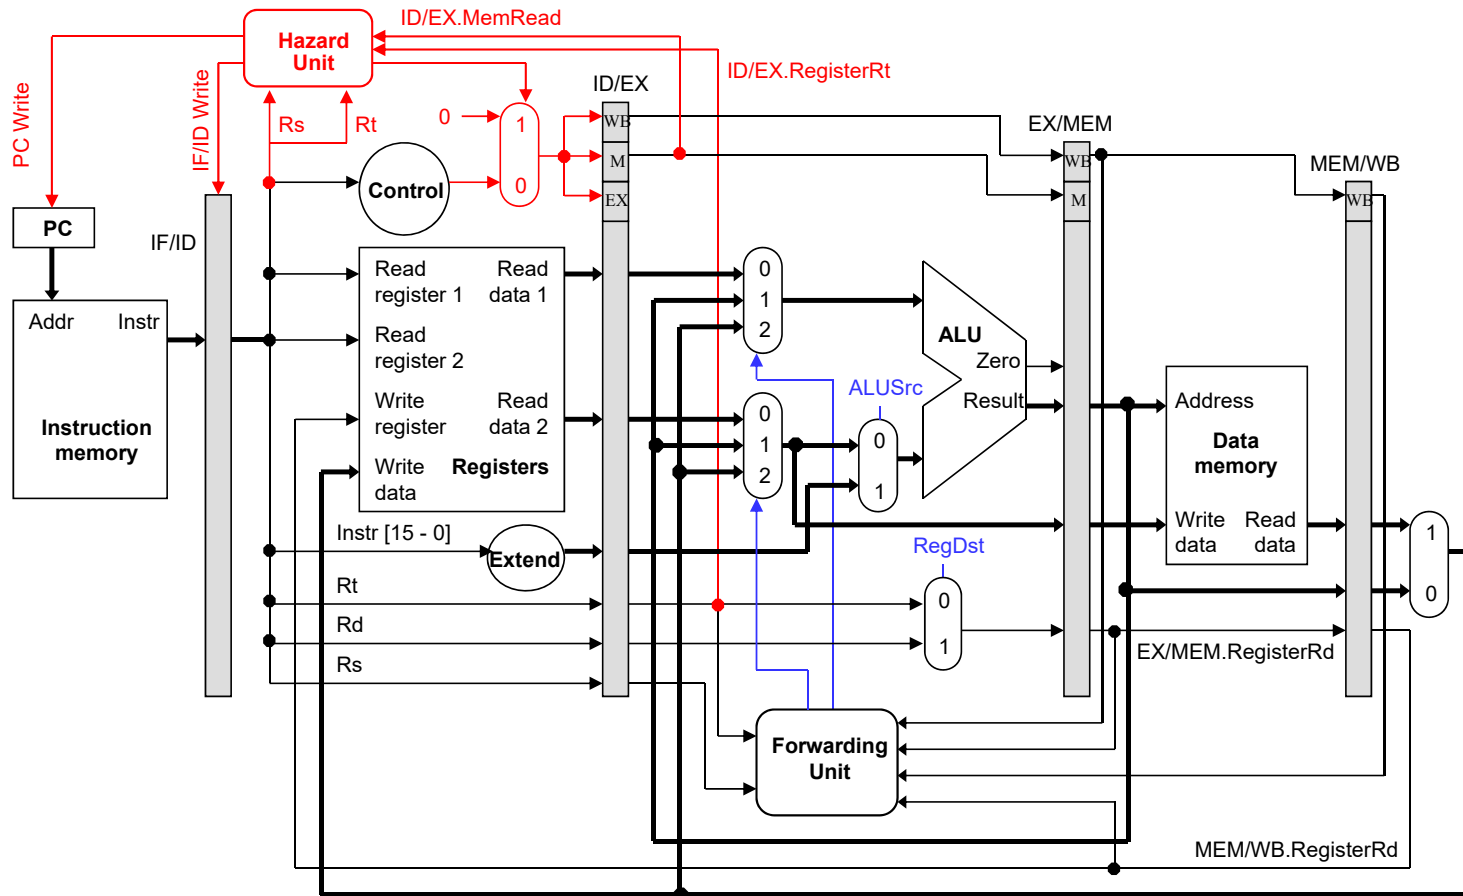
	1	2	3	4	5	6	7	8	9	10	11
lw	F	D	E	M	W						
lw		F	D	D	E	M	W				
and			F	F	D	D	E	M	W		
add					F	F	D	E	M	W	
						A	B	C	D	E	

In which cycle is
 add \$4, \$9, \$2
 in the WB stage?

Adding hazard detection to the CPU



Adding hazard detection to the CPU

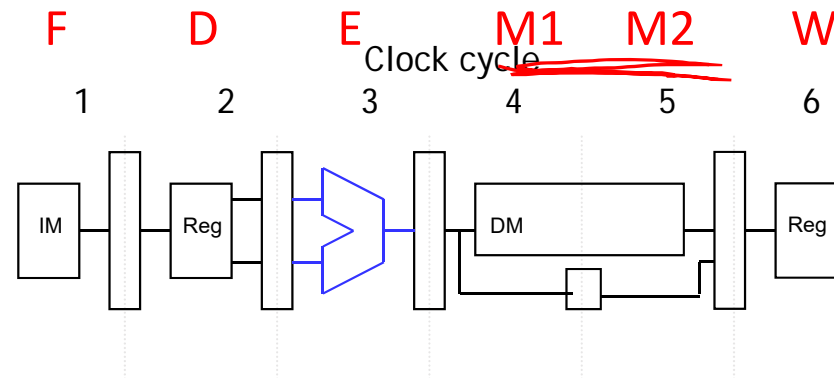


The hazard detection unit

- The hazard detection unit's inputs are as follows.
 - `IF/ID.RegisterRs` and `IF/ID.RegisterRt`, the source registers for the current instruction.
 - `ID/EX.MemRead` and `ID/EX.RegisterRt`, to determine if the previous instruction is LW and, if so, which register it will write to.
- By inspecting these values, the detection unit generates three outputs.
 - Two new control signals `PCWrite` and `IF/ID Write`, which determine whether the pipeline stalls or continues.
 - A `mux select` for a new multiplexer, which forces control signals for the current EX and future MEM/WB stages to 0 in case of a stall.

Generalizing Forwarding/Stalling

- What if data memory access was so slow, we wanted to pipeline it over 2 cycles?

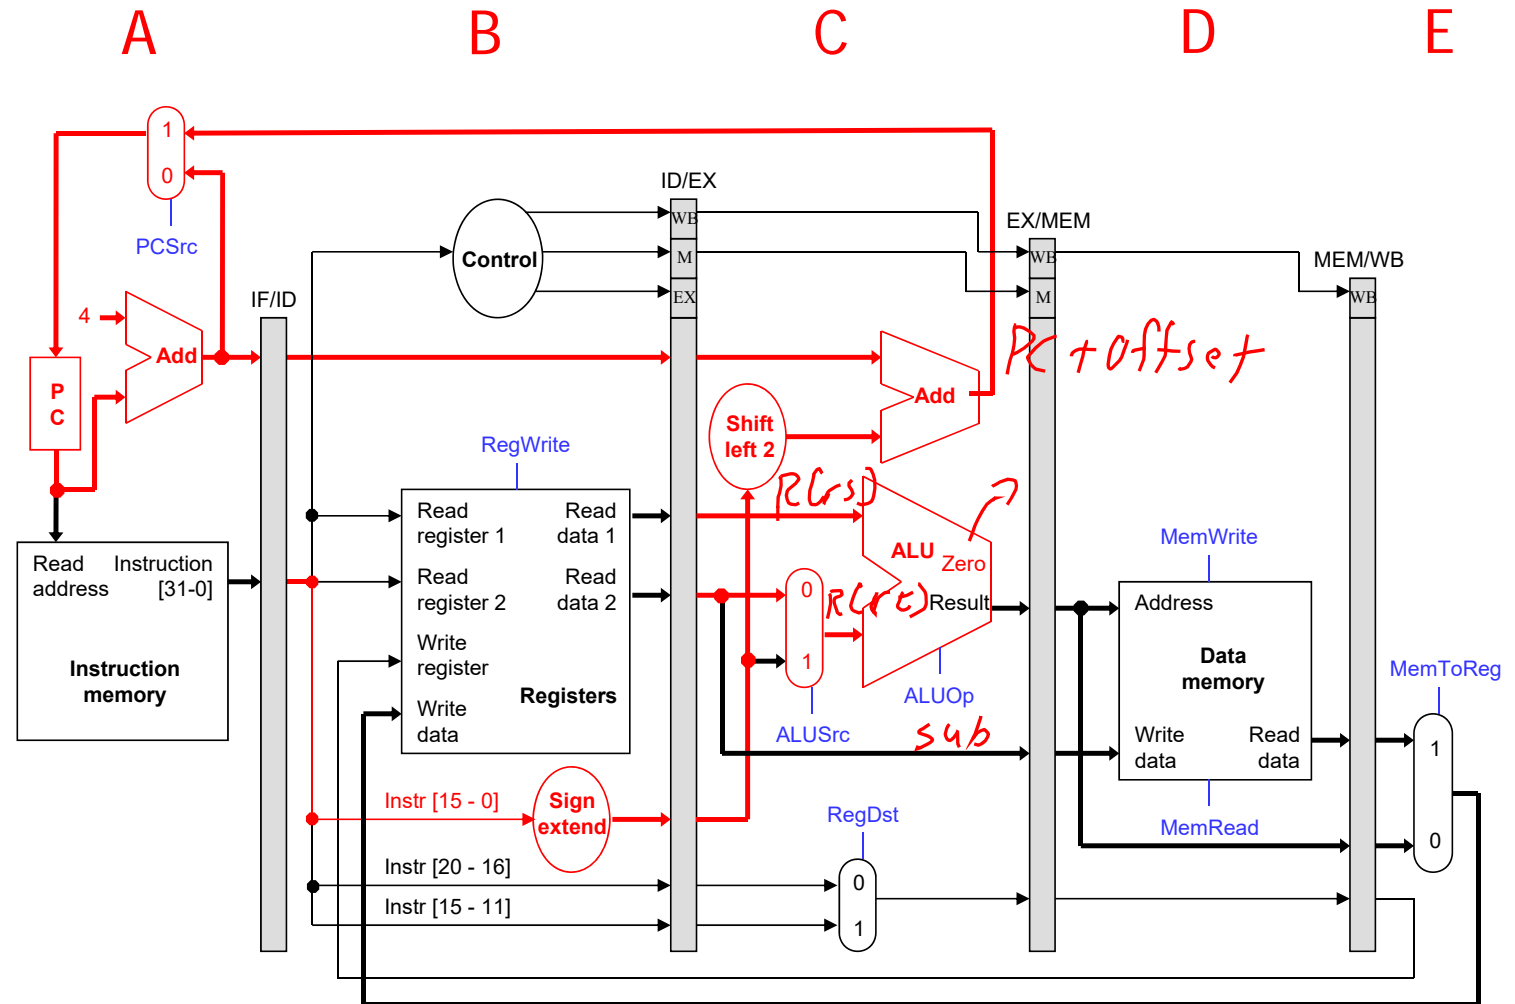


- How many bypass inputs would the muxes in EXE have?
- Which instructions in the following require stalling and/or bypassing?

		1	2	3	4	5	6	7	8	9	10	11	12	13
lw	\$13, 0(\$11)	F	D	E	M1	M2	W							
add	\$7, \$8, \$9		F	D	E	M1	M2	W						
add	\$15, \$7, \$13			F	D	D	E	M1	M2	W				
								A	B	C	D	E		

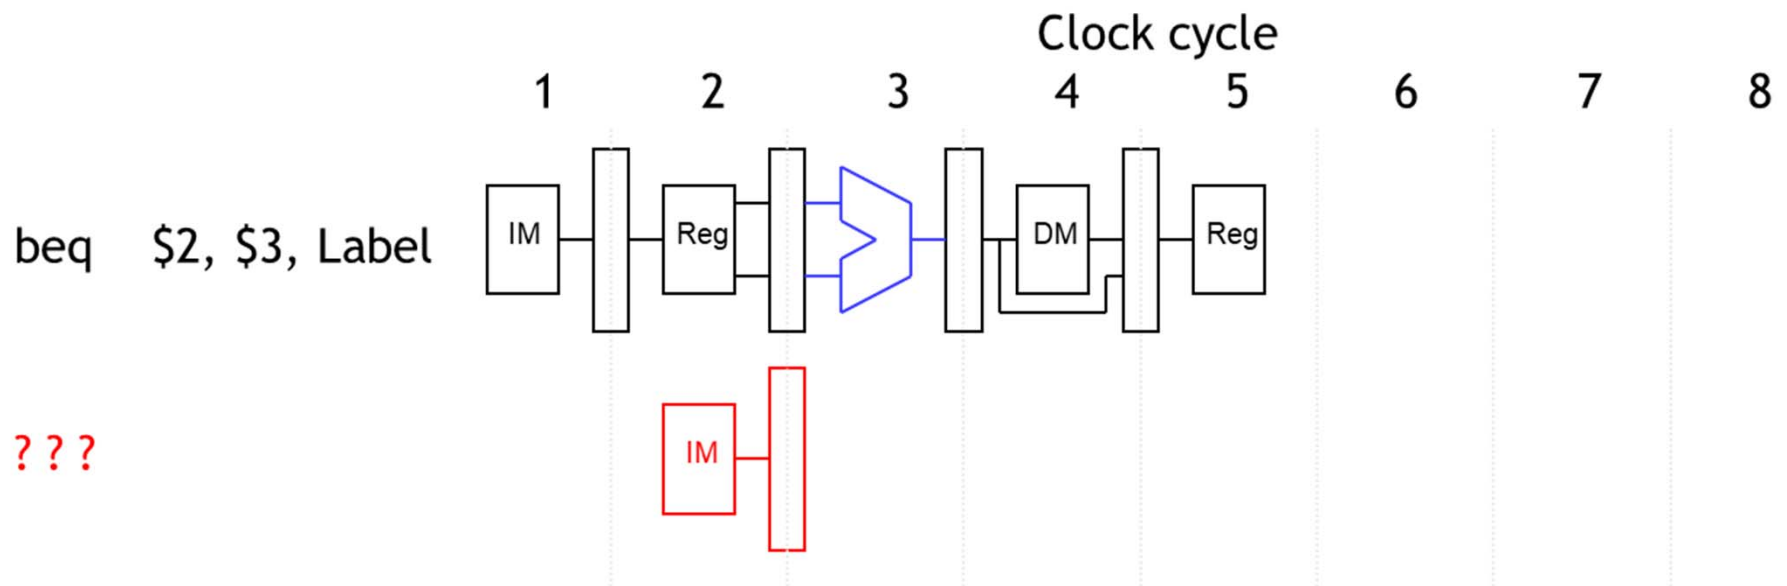
In which cycle is add \$15, \$7, \$13 in the WB stage?

When are branches resolved (when do I know for certain where I am branching to)?

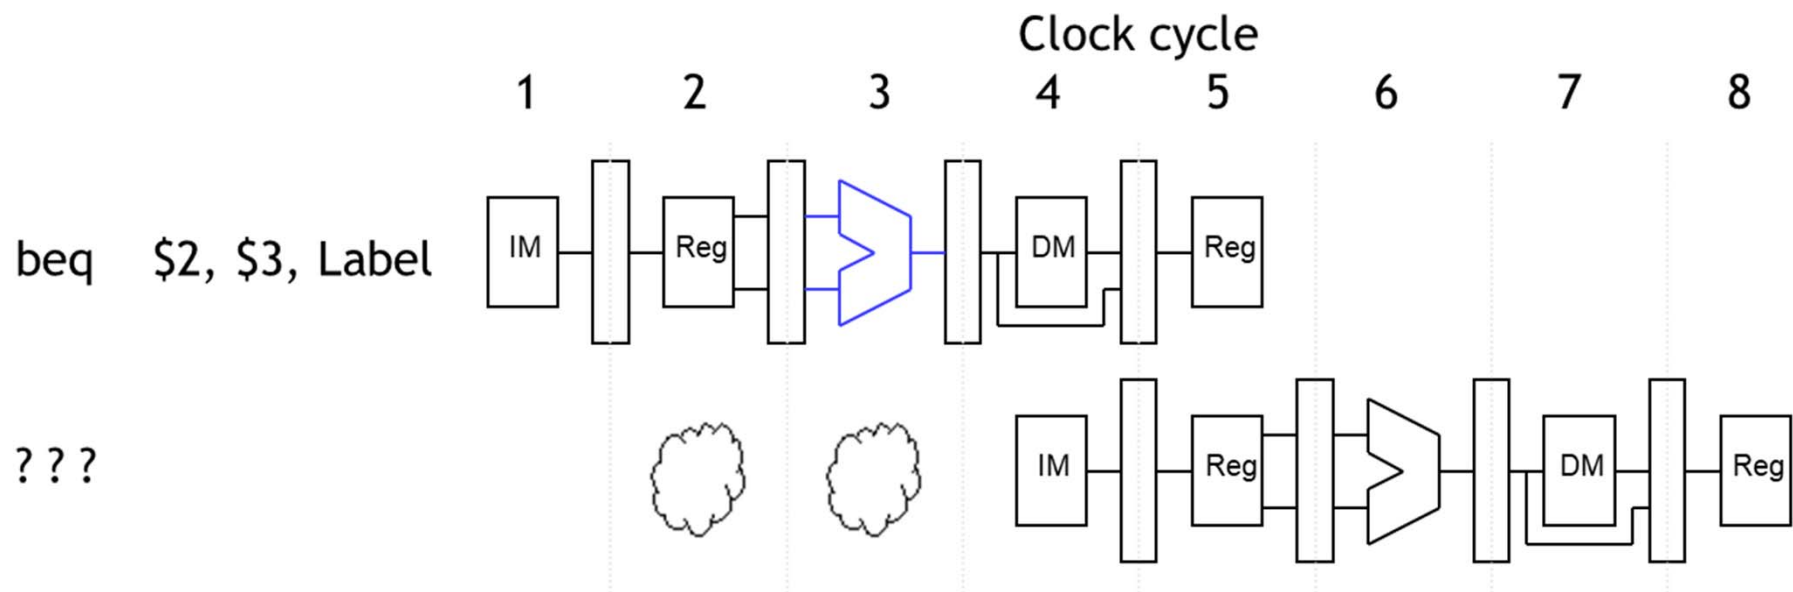


Branch target address and zero flag are determined in the EX stage, creating a **control hazard**

- Therefore, branch decision cannot be made until the end of the EX stage.
- But we need to know which instruction to fetch next, in order to keep the pipeline running!



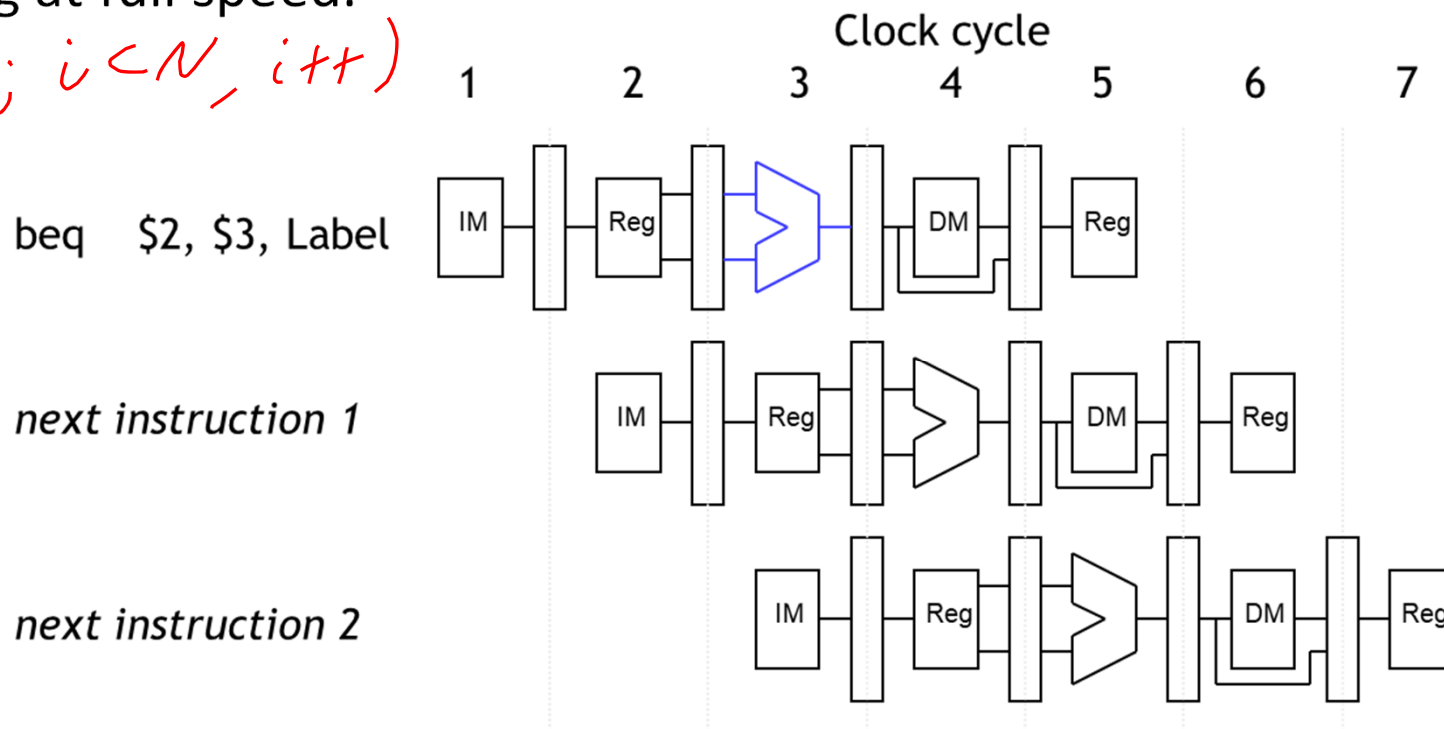
Could handle the hazard by stalling until cycle 4 after the branch destination is known



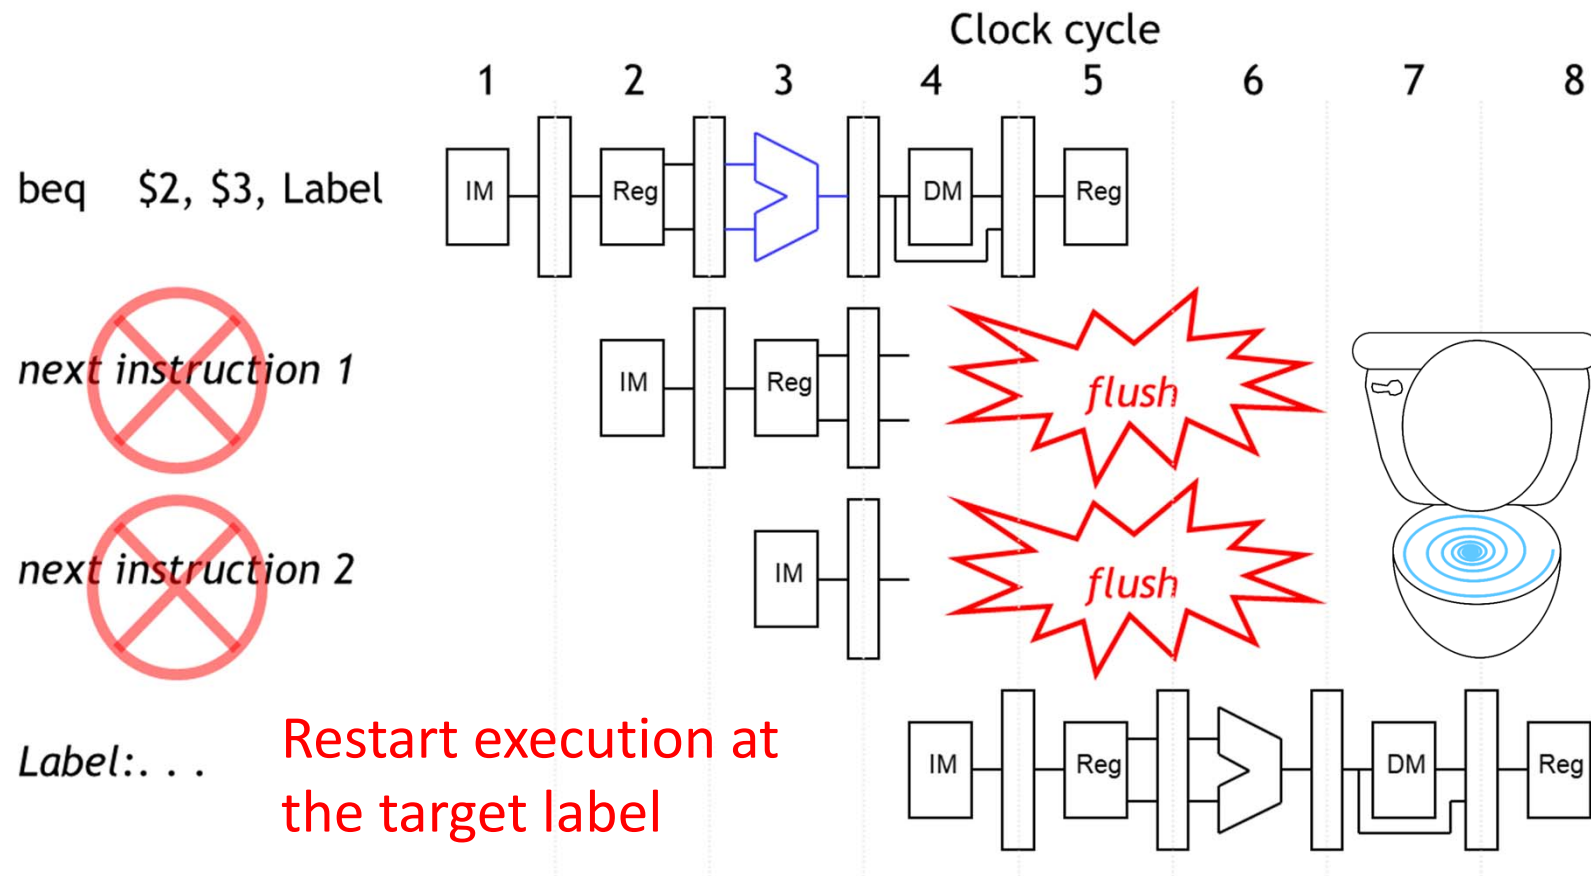
Alternatively, we could guess whether the branch is taken (i.e., **Branch Prediction**)

- Easy to just assume the branch is *not* taken, just increment PC
- If we're correct, then there is no problem and the pipeline keeps going at full speed.

for (i=0; i<N; i++)



If we mispredict the branch, we need to **flush** two incorrect instructions from the pipeline



i>clicker

Which method of handling branches will generally lead to better performance?

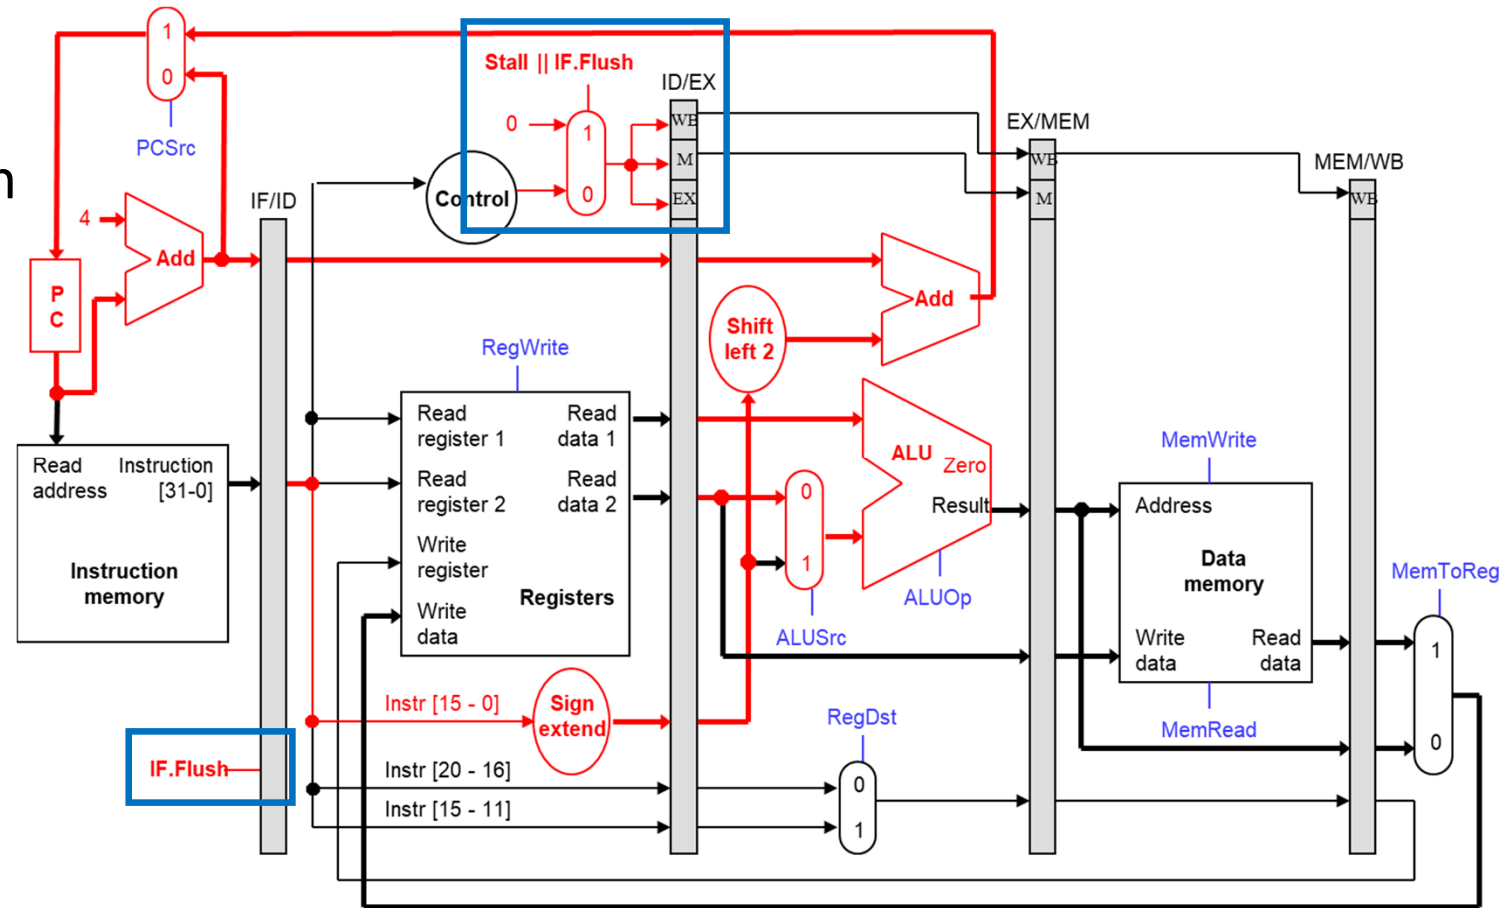
- a) Stalling
- b) Branch prediction *+ flushes*
- c) They will give about the same performance

On a misprediction, branch prediction wastes two cycles; stalling always wastes those two cycles

- All modern CPUs use branch prediction.
 - Accurate predictions are important for optimal performance.
 - Most CPUs predict branches dynamically—statistics are kept at run-time to determine the likelihood of a branch being taken.
- The pipeline structure also has a big impact on branch prediction.
 - A longer pipeline may require more instructions to be flushed for a misprediction, resulting in more wasted time and lower performance.
 - We must also be careful that instructions do not modify registers or memory before they get flushed.

Flush the pipeline by turning the instructions in ID/EX and IF/ID into nop

A flush is required when the instruction in EX is a BEQ and “zero” is 1.



Summary

- Three kinds of hazards conspire to make pipelining difficult.
- **Structural hazards** result from not having enough hardware available to execute multiple instructions simultaneously.
 - These are avoided by adding more functional units (e.g., more adders or memories) or by redesigning the pipeline stages.
- **Data hazards** can occur when instructions need to access registers that haven't been updated yet.
 - Hazards from R-type instructions can be avoided with forwarding.
 - Loads can result in a “true” hazard, which must stall the pipeline.
- **Control hazards** arise when the CPU cannot determine which instruction to fetch next.
 - We can minimize delays by doing branch tests earlier in the pipeline.
 - We can also take a chance and predict the branch direction, to make the most of a bad situation.