

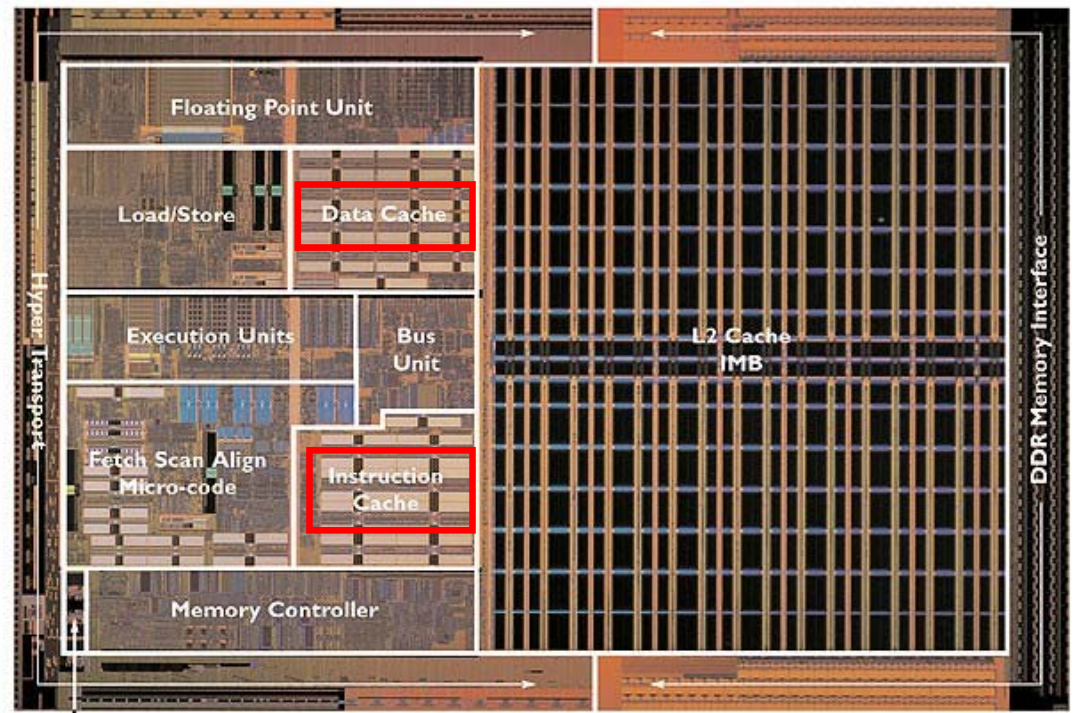
Cache performance and memory access patterns

Today's lecture

- Trade-offs between cache-size, block-size, associativity and performance (miss rate)
- Predicting Cache Performance from Code
 - Memory access patterns
 - Cache hit analysis for a given address stream


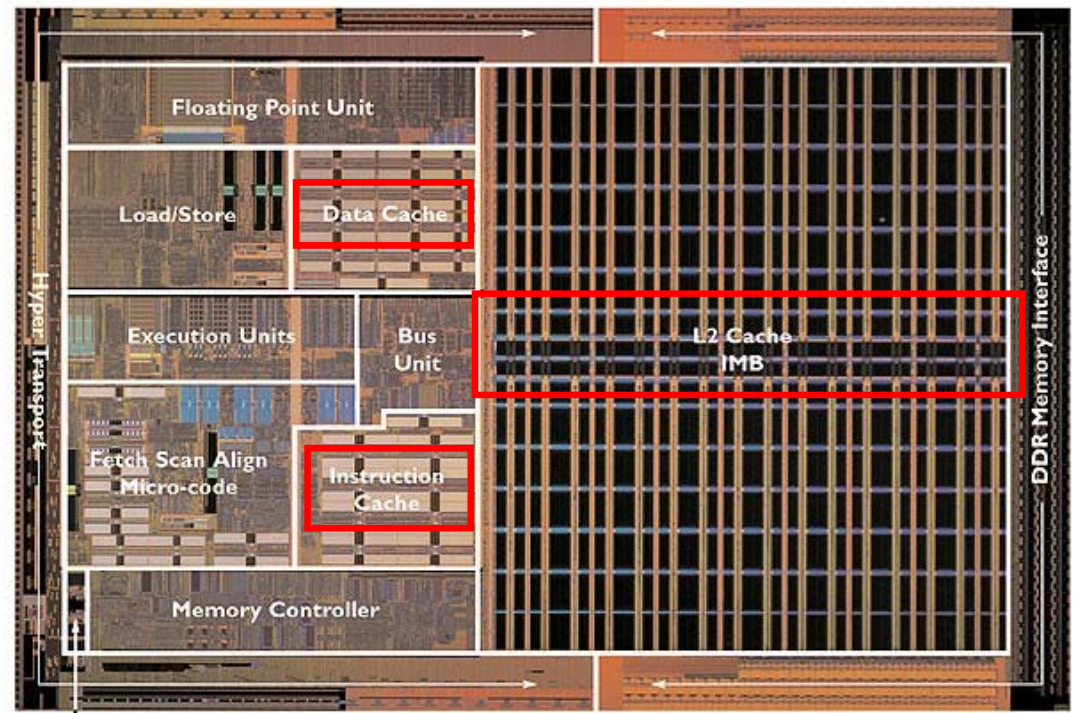
Most real architectures split instruction and data caches

- Pro: No structural hazard between IF & MEM
 - A single-ported, unified cache stalls fetch during load or store
- Con: Static partitioning of cache between instructions & data
 - Bad if working sets unequal: *e.g.*, code/**DATA** or **CODE**/_{data}




Most real architectures use multiple levels of caches

- Trade-off between access time & hit rate
 - L1 cache focuses on fast access time (with okay hit rate)
 - L2 cache focuses on good hit rate (with okay access time)



Example cache statistics from Opteron



- L1 Caches: Instruction & Data
 - 64 kB
 - 64 byte blocks
 - 2-way set associative
 - 2 cycle access time
- L2 Cache:
 - 1 MB
 - 64 byte blocks
 - 4-way set associative
 - 16 cycle access time (total, not just miss penalty)
- Memory
 - 200+ cycle access time

Tradeoffs between cache parameters are evaluated experimentally

- General goal is to minimize miss rate
- Will look at examples on the following slides
- We will do some cache simulations on the MP's.

Remember: miss rate is the number of misses per memory access


Predict associativity vs. miss rate

Increasing associativity will

- a) Decrease the miss rate
- b) Increase the miss rate
- c) Have mixed results

Increasing associativity generally improves miss rates at the cost of hardware complexity

- Each set has more blocks, so there's less chance of a conflict between two addresses which both belong in the same set.




Predict cache size vs. miss rate

Increasing cache size will

- a) Decrease the miss rate
- b) Increase the miss rate
- c) Have mixed results

Increasing cache size also generally improves miss rates

- The larger a cache is, the lower chance of conflict.




Predict block size vs. miss rate

Increasing block size will

- a) Decrease the miss rate
- b) Increase the miss rate
- c) Have mixed results

Small blocks do not take advantage of spatial locality

- But if cache blocks get too big relative to the cache size, they increase conflicts



Cache Size Relationships

- $\text{NUM_BLOCKS} = \text{NUM_SETS} * \text{NUM_BLOCKS/SET}$
- $\text{CACHE_SIZE} = \text{NUM_BLOCKS} * \text{BLOCK_SIZE}$


- Example:

- 2-way set associative, 512 sets, 32B blocks

- $\text{NUM_BLOCKS} =$ a) 32 b) 64 c) 512 d) 1024

- $\text{CACHE_SIZE} =$
 - a) 32B
 - b) 512B
 - c) 1KB
 - d) 16KB
 - e) 32KB

$1k = 1024$



Code -> Address stream (Example)

```
int A[SIZE], total = 0;
for (int i = 0 ; i < SIZE ; i ++ ) {
    total += A[i];
}
```

lw A[i]

$total = total + A[i]$

- How many loads/stores are there in this piece of code?
- What is the stream of addresses generated?

assume $A = \&A[0]$

$A, A+4, A+8, A+12 \dots$

We can estimate the miss rate based on the address stream

- Address stream: A, A+4, A+8, A+12, A+16, A+20, A+24, A+28, A+32, A+36, A+40, A+44, A+48, A+52, A+56, A+60, A+64, A+68, ...
- Direct-mapped cache, 1024 x 32B blocks (total size: 32,768B)
- What is the miss rate if the stream is smaller than the cache (i.e., SIZE is smaller than 1032 x 32)?


1 miss / 8 accesses
1/8 miss rate

M	H	H	H	H	H	H	H
A[0]	A[1]	2	Cache	4	5	6	7
A[8]	A[9]	A[10]	11	12	13	14	15
16	17	18	19	20	21	22	23
M	H	H	H	H	H	H	H

...

i>clicker question

- Address stream: A, A+4, A+8, A+12, A+16, A+20, A+24, A+28, A+32, A+36, A+40, A+44, A+48, A+52, A+56, A+60, A+64, A+68, ...
 - Direct-mapped cache, 1024 x 32B blocks (total size: 32,768B)
 - What is the miss rate if the stream is greater than the cache?
- a) $>1/8$ b) $1/8$ c) $<1/8$



Adjacent memory blocks go to adjacent cache sets

- Assume that $A = 0x00010038$

A	$=$	$0x00010038$	$=$	0001	0000	0000	0011	1000
$A+4$	$=$	$0x0001003c$	$=$	0001	0000	0000	0011	1100
$A+8$	$=$	$0x00010040$	$=$	0001	0000	0000	0100	0000
$A+12$	$=$	$0x00010044$	$=$	0001	0000	0000	0100	0100
...								
$A+40$	$=$	$0x0001005c$	$=$	0001	0000	0000	0101	1100
$A+44$	$=$	$0x00010060$	$=$	0001	0000	0000	0110	0000

1024 x 32B Blocks -> 5 bit block offset, 10-bit set index

Estimate the miss rate for a two-way set-associative cache with the same number of blocks

- Address stream: A, A+4, A+8, A+12, A+16, A+20, A+24, A+28, A+32, A+36, A+40, A+44, A+48, A+52, A+56, A+60, A+64, A+68, ...
- Cache: Two-way set associative, 512 sets, 32B blocks (same size)



LRU	M	11	11	11	11	11	11	11
LRU	0	1	2	3	4	5	6	7
LRU	M	11	11	11	11	11	11	11
LRU	M							
LRU	M	11	11	11	11	11	11	11

LRU	M							
513	M	11	11	11	11	11	11	11
8	M	9	10	11	12	13	14	15
LRU	M	11	11	11	11	11	11	11
LRU	M							

a) $> 1/8$ b) $1/8$ miss rate c) $< 1/8$

Stride length is the amount we increment over an array

```
int A[SIZE], total = 0;
for (int i = 0 ; i < SIZE ; i += STRIDE) {
    total += A[i];
}
```

// STRIDE != 1

- If STRIDE = 2, what is the stream of addresses generated?

A, A+8, A+16, A+24

What is the miss rate for a stride length of 2?

```
int A[SIZE], total = 0;
for (int i = 0 ; i < SIZE ; i += STRIDE2) { // STRIDE2
    total += A[i];
}
```

Cache: 2-way SA, 512 sets, 32B blocks

- a) $1/32$ b) $1/16$ c) $1/8$ d) $1/4$ e) Not enough info

LRU

M 14		H 16		H 18		H 20	
M 22		H 24		H 26		H 28	

LRU

LRU

LRU

		M A[0]		H A[2]		H A[4]	
		H A[6]		H A[8]		H A[10]	

LRU


LRU

Strided access generally increases the miss rate because it reduces spatial locality


```
int A[SIZE], total = 0;
for (int i = 0 ; i < SIZE ; i += STRIDE) { // STRIDE != 1
    total += A[i];
}
```

In general, two scenarios:

STRIDE < CACHE BLOCK SIZE



STRIDE >= CACHE BLOCK SIZE



The miss rate of strided access can be expressed mathematically

- Miss Rate = $\text{MIN}(1, \text{DATA_SIZE} * \text{STRIDE} / \text{CACHE_BLOCK_SIZE})$

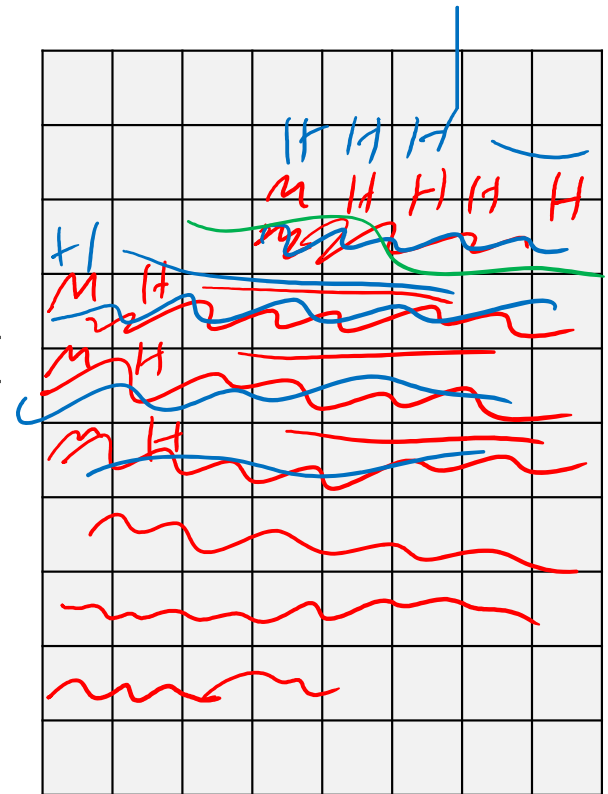
Nested loops create temporal locality in data access

```
int A[SIZE], total = 0;
for (int j = 0 ; j < N ; j ++ ) {
    for (int i = 0 ; i < SIZE ; i ++ ) {
        total += A[i];
    }
}
```

What does stream of addresses look like?

$A, A+4, A+8, A+12, \dots$

$A, A+4, A+8$



If $\text{sizeof(int)} * \text{SIZE} < \text{CACHE_SIZE}$ and N is large, what is the miss rate?

- a) 1/32 b) 1/2 c) 1 d) 0 e) not enough info

iter 1: $\frac{1}{8}$ miss rate
 iter 2: 0 miss rate

(pick best answer)

If the data structure is too big, we lose temporal locality

```
int A[SIZE], total = 0;
for (int j = 0 ; j < N ; j ++ ) {
    for (int i = 0 ; i < SIZE ; i ++ ) {
        total += A[i];
    }
}
```

What does stream of addresses look like?




If $\text{sizeof}(\text{int}) * \text{SIZE} \geq 2 * \text{CACHE_SIZE}$ and N is large, what is the miss rate?

$\sim \frac{1}{8}$ miss rate


iter 1: $\frac{1}{8}$ miss rate

iter 22: $\frac{1}{8}$ miss rate





Performance depends on both the cache and data structure properties

```
int A[SIZE], total = 0;
for (int j = 0 ; j < N ; j ++ ) {
    for (int i = 0 ; i < SIZE ; i ++ ) {
        total += A[i];
    }
}
```



$\text{CACHE_SIZE} \leq \text{sizeof(int)} * \text{SIZE} \leq 2 * \text{CACHE_SIZE}$



Non-linear access of arrays sacrifices temporal and spatial locality

```
int A[SIZE], total = 0;
for (int j = 0 ; j < 4 ; j ++ ) {
    for (int i = 0 ; i < SIZE ; i += 4) {
        total += A[i+j];
    }
}
```

What does stream of addresses look like?

$A[0], A[4], A[8], A[12], \dots$
 $A[1], A[5], A[9], A[13], \dots$

What is the miss rate?

$$\frac{1}{2}$$

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15

$$\frac{1}{8}$$

General Reuse Questions


- Is the data still in the cache when reuse occurs?

Temporal Locality

- Special case: Does all of the data fit in the cache?

i>clicker, how do linked lists perform on a cache?

```
struct list_t { int data, struct list_t *next };  
for (struct list_t *l = list_head ; l != NULL ; l = l->next) {  
    total += l->data;  
}
```




Assume a 2-way SA cache with 512 sets, 32B blocks.

If the list is 1000 elements long, how many misses per iteration of the loop?

- a) 1 b) $\frac{1}{4}$ c) $\frac{1}{8}$ d) 0 e) not enough info

i>clicker Multiple Data Structures

```
int A[SIZE], B[SIZE];  
for (int i = 0 ; i < SIZE ; i ++ ) {  
    B[i] = A[i];  
}
```



What is the stream of addresses generated?

Assume a direct-mapped cache with 1024 32B blocks.

Assume $\text{sizeof}(\text{int}) * \text{SIZE} \gg \text{CACHE_SIZE}$.

How many cache misses per iteration of the loop?

- a) 2 b) 1 c) 1/4 d) 1/8 e) not enough info

What addresses map to the same set?

- Multiples of the set size.
- $\text{SET_SIZE} = \text{NUM_SETS} * \text{BLOCK_SIZE}$
 - Also $\text{SET_SIZE} = \text{CACHE_SIZE} / \text{NUM_BLOCKS/SET}$
- Examples:
 - Direct-mapped, 1024 32B blocks
 - 2-way SA, 512 sets, 32B blocks

Cache Analysis Questions

- How big is each data item? (how many fit in a cache block?)
- Is there data reuse or use of data in the same cache block?
- Is the data/block still in cache when reuse occurs?
 - **How big is the data relative to the cache?**
 - **Is there aliasing (cache conflicts) that is problematic?**
- Some useful relationships:
 - Strides: $\text{miss rate} = \text{MIN}(1, \text{sizeof}(\text{data}) * \text{stride} / \text{BLOCK_SIZE})$
 - Adjacent blocks in memory \rightarrow adjacent sets
 - Addresses $A, A + \text{BLOCK_SIZE} * \text{NUM_SETS}$ are in the same set