

# CS 398 ACC

## MapReduce Part 2

---

Prof. Robert J. Brunner

Ben Congdon

Tyler Kim

# Change in Quiz Policy

Starting with Quiz 2:

- Only **two** attempts allowed
  - But, you can see all that you got right/wrong

MP1

How's it going?

# Tentative Auto-Grader Schedule

- Wednesday Evening (~9pm)
  - Friday Evening (~9pm)
  - Sunday Evening (~9pm)
  - Monday Evening (~9pm)
  - Tuesday Midday (~2pm)
- 
- Results location: /mp1/grade\_report.txt
  - This will be posted on Piazza as well.

# Outline

- MapReduce Programming
  - Word Count Implementation
  - Conventions / Pitfalls
- Execution Options
- MapReduce Use Cases

# Outline

- **MapReduce Programming**
  - Word Count Implementation
  - Conventions / Pitfalls
- Execution Options
- MapReduce Use Cases

# Reminders...

$\langle \text{key\_input}, \text{val\_input} \rangle$	$\Rightarrow$	$\langle \text{key\_inter}, \text{val\_inter} \rangle$	$\Rightarrow$	$\langle \text{key\_out}, \text{val\_out} \rangle$
	<i>Map</i>		<i>Reduce</i>	

- **Map:**

- **Input:** “Original” input data or key/value pairs from previous chained job
- **Output:** Intermediate key/value pairs

- **Reduce:**

- **Input:** Intermediate key/value pair (per key)
- **Output:** Final key/value pairs

# Word Count - Mapper

...

```
class WordCount(MRJob):  
    def mapper(self, key, val):  
        for word in WORD_REGEX.findall(val):  
            yield (word, 1)
```

...



# Word Count - Reducer

```
def reducer(self, key, vals):  
    total_sum = 0  
  
    # Iterate and count all occurrences of the word  
    for v in vals:  
        total_sum += 1  
  
    # Yield the word and number of occurrences  
    yield key, total_sum
```

# Reminders...

- Python Iterators

```
def more_efficient(self, key, values):  
    for v in values:  
        yield key, v + 10
```

```
def inefficient(self, key, values):  
    # List comprehension loads all values into memory  
    plus_10 = [v + 10 for v in values]  
    for v in plus_10:  
        yield key, v
```

# MapReduce - Common Conventions

- **Composite Key**

- Use more than one attribute in the construction of a key
  - i.e. <(city, state), population>

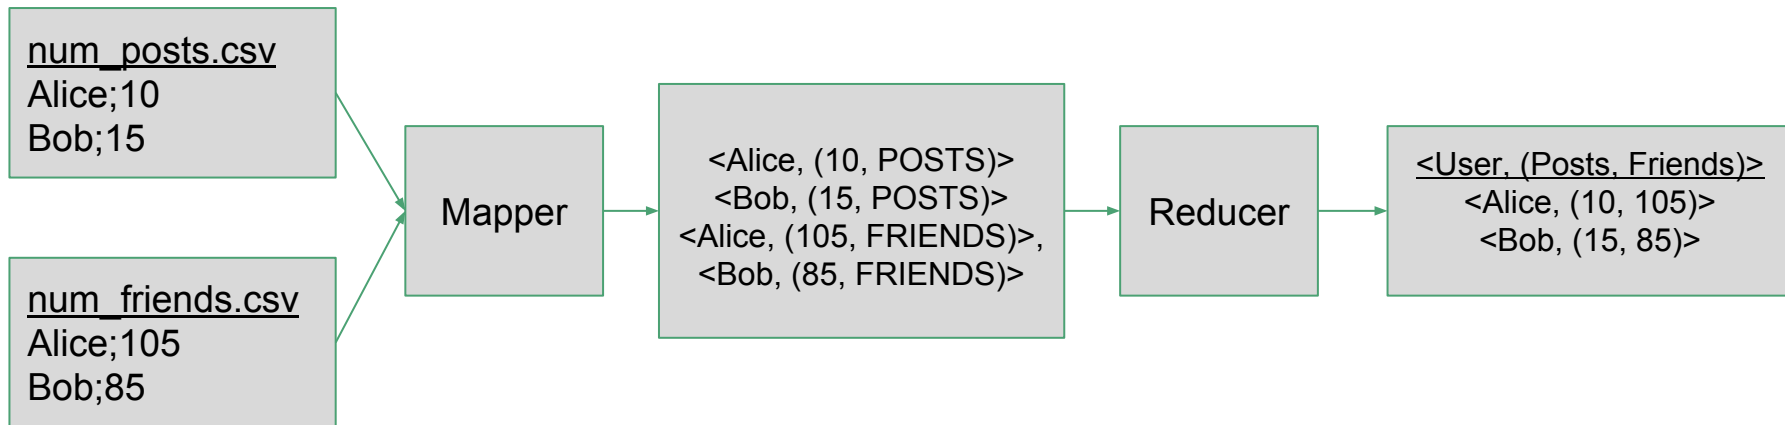
- **Composite Value**

- Use more than one attribute in the construction of a value
  - i.e. <user\_id, (num\_friends, num\_posts)>
- Can also use custom serialization methods for intermediate values
  - i.e. JSON, Python Pickling
  - Just be careful about size overhead (for bandwidth)

# MapReduce - Common Conventions

- **Joins**

- Idea: Use input datasets with more than one format
- Mapper: Add flag to output to indicate value type
- Reducer: Reconcile attributes by key into a single record



# MapReduce - Common Pitfalls

- **“Leaky” Reducers**

- **Source:** Reducers that use too much memory (i.e. keeping all values in memory)
  - Reducing functions have “too much” state
  - Might not be due to bad reducer design, but rather empirical workload/machine limitations

# MapReduce - Common Pitfalls

- “Leaky” Reducers

- Example:

```
# Want to count number of unique values

def reduce(self, key, values):
    # 'unique' will store all distinct values we see
    unique = set()
    for v in values:
        if v not in unique:
            unique.add(v)

    # yield size of unique value set
    yield (key, len(unique))
```

# MapReduce - Common Pitfalls

- **“Leaky” Reducers**

- **Source:** Reducers that use too much memory (i.e. keeping all values in memory)
  - Reducing functions have “too much” state
  - Might not be due to bad reducer design, but rather empirical workload/machine limitations
- **Solutions:**
  - Benchmark workload - Maybe it’s not an issue
  - Take advantage of secondary sort
  - Use the fact that values are passed as an iterator (Use a “stream” mindset)

# MapReduce - Common Pitfalls

- “Leaky” Reducers

- Example (Fixed):

```
# Want to count number of unique values

def reduce(self, key, values):
    prev, count = None, 0

    # Here we assume that values is sorted
    for v in values:
        if v != prev:
            count += 1
            prev = v

    yield (key, count)
```



# MapReduce - Common Pitfalls

- **“Hot” Keys**
  - **Source:** Some keys may contain many, many more values than most other keys

# MapReduce - Common Pitfalls

- “Hot” Keys

- Example:

- Google Web Indexing
    - Average Key Size: 300 KB
    - Some keys have 50+ GB

Reduce key	Reduce input size
*.blogspot.com	82.9G
cgi.ebay.com	58.2G
profile.myspace.com	56.3G
yellowpages.superpages.com	49.6G
www.amazon.co.uk	41.7G
average reduce input size for a given key	300K

# MapReduce - Common Pitfalls

- **“Hot” Keys**

- **Source:** Some keys may contain many, many more values than most other keys
- **Solutions:**
  - Benchmark workload - Maybe it's not an issue
  - Write a custom partitioner so that load is still distributed evenly across machines
    - Partitioner determines which intermediate keys go to which reducer
    - Goal: Distribute load evenly so “hot keys” don't all go to one reducer
  - Can you split up keys further and recombine them in a chained MR step?

# MapReduce - Common Pitfalls

- Python Specific: **Handling Types**
  - Input data is almost always ``str` / `bytes``
  - Java allows you to define custom types and serialization
  - Python can do this too (i.e. with Pickling), but it is not required

# Outline

- MapReduce Programming
  - Word Count Implementation
  - Conventions / Pitfalls
- **Execution Options**
- MapReduce Use Cases

# MapReduce Execution Options

- Hadoop (native Java)
- Hadoop Streaming
- External Frameworks (MRJob)

# MapReduce Programming Options

- **Hadoop MapReduce (native)**

- Preferred, most performant method for writing Hadoop MapReduce jobs
- Minimum Required Components per Job:
  - Mapper, Reducer, Job execution boilerplate
- Additional Customizable Components:
  - **InputFormat**: Splits input files into chunks to be distributed to mappers
  - **Partitioner**: Controls which keys go to which reducers (default: HashPartitioner)
  - **OutputFormat / OutputComitter**: Handles the end of the reduce phase (usually writing job output to disk)

# MapReduce Programming Options

- **Hadoop Streaming**

- Hadoop copies arbitrary binary executable(s) for mappers and reducers
- Uses STDIN/STDOUT to stream data to mappers/reducers
  - **Mapper:** Each input record is a new line
  - **Reducer:** You receive a stream of arbitrary k/v pairs (sorted by key)
    - You (the program) have to figure out when you switch from one key to the next
- Still parallel, but difficult to work in
  - Everything is text; key/values are usually just tab separated
- Allows non-Java languages to be used on the Hadoop Framework



# MapReduce Programming Options

- Aside: Local Unix Commands
  - `$ echo $DATA | ./mapper | sort -k1,1 | ./reducer > output`
  - Not parallel (or advisable), but useful for debugging Hadoop Streaming jobs
  - Similar to Hadoop Streaming in that you use an arbitrary executable as a mapper / reducer

# MapReduce Programming Options

- **MRJob (External Framework)**
  - Python Framework for writing / running MapReduce jobs
  - Built by Yelp
  - Write Once, Run “Anywhere” (actually, though!)
  - Supported Execution Environments:
    - Local Execution (MP1)
    - Hadoop (MP2)
    - GCP Dataproc - Google’s Hosted MapReduce
    - AWS EMR - Amazon’s Hosted Hadoop

# MapReduce Programming Options

- **MRJob (External Framework)**
  - How does it work?
    - Hadoop Streaming under-the-hood
    - Provides similar abstractions as the Native Java API

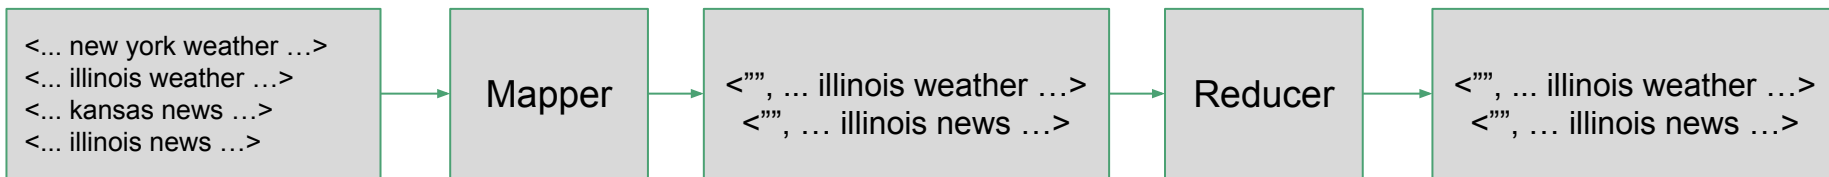
# Outline

- MapReduce Programming
  - Word Count Implementation
  - Conventions / Pitfalls
- Execution Options
- **MapReduce Use Cases**

# Distributed Grep

Used for: Filtering, Parsing, or Validation

- **Input:** Large set of files
- **Mapper:** Look at input and emit records containing the query term
- **Reducer:** Pass through all records unchanged

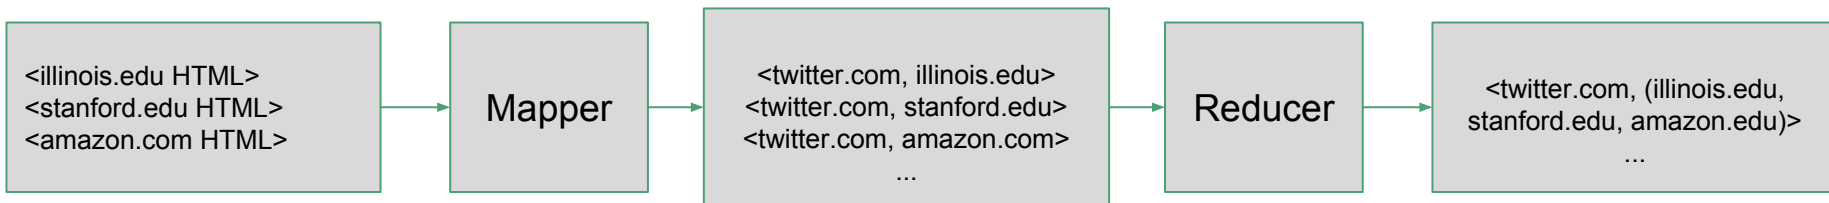


# Graph Processing

Web-Linked / Web Scraping Graph (Similar to Problem 2 of MP1)



- **Input:** HTML Text
- **Map output:** <target, source> pairs
  - i.e. Search for <a href="...">
- **Reduce output:** <target, list(source)> pairs

Used by Google for web search indexing



# Geospatial / Satellite Data

- **Input:** Geospatial coordinates, satellite data
- **Map output:**
  - `<map_tile, tile_information>` pairs
  - “Chunk” geographic region by tile
- **Reduce output:**
  - `<map_tile, final_tile_info>` pairs



Used by Google Maps to reconcile satellite imagery over time

# What is required:

## **Programmer:**

1. Don't need to know specifics about parallel/distributed computing/programming.
2. Know the data source/format
3. Write a map/reduce programs
4. Submit jobs and wait :)




# What is required:

**Framework/Library** (e.g. mrjob, hadoop, etc.):

1. Parallelize Map (distribute to mapping machines)
2. Transfer Data / Shuffle Data
3. Parallelize Reduce (distribute to reducing machines)
4. Deal with failure, missing values
5. Implement data transfer. Input/Output/Artifacts. Interact with distributed file system

## Next Week:



- Lets you run MapReduce on many computers for a single task.
- Can scales to 1000s of nodes
- Processes Petabytes with (relative) ease
- Get on course cluster!